

Advanced Kubernetes

=====

Lab 2 – Kubelet and advanced pod specification

The `kubelet` is the primary "node agent" that runs on each Node in a Kubernetes cluster. The `kubelet` works in terms of PodSpecs. A PodSpec is a YAML or JSON object that describes a pod. The `kubelet` takes a set of PodSpecs that are provided through various mechanisms (primarily through the kube-apiserver), and ensures that the containers described in those PodSpecs are running and healthy.

Other than a PodSpec from the kube-apiserver, there are three additional ways that a pod manifest can be provided to the `kubelet` :

- **File** - The `--pod-manifest-path` switch can be used to pass a path containing pods to run on startup. This path is rechecked every 20 seconds (configurable with a flag)
- **HTTP endpoint** - HTTP endpoint passed as a parameter on the command line, checked every 20 seconds (configurable)
- **HTTP server** - The `kubelet` can also listen for HTTP manifest posts

The `--pod-manifest-path` is typically used to tell the kubelet to start other kubernetes components, like the kube-proxy on worker nodes and the kube-apiserver on master nodes.

1. Stop running cluster components

To experiment with the `kubelet` independently, stop the Kubernetes components (kube-apiserver, kubelet, and etcd) you may have running by typing ^C (or kill - SIGINT) in their TTYs.

```
user@nodea:~$ sudo kill -SIGINT $(pidof kube-apiserver kubelet etcd)

user@nodea:~$
```

Before you restart your cluster, clear the kube-apiserver cluster state by removing the `etcd` backing store:

```
user@nodea:~$ rm -Rf ~/default.etcd/
user@nodea:~$
```

The `kubelet` also caches its state on disk. You can eliminate the kubelet's cached state by stopping kubelet and then removing the kubelet's backing store as well:

```
user@nodea:~$ sudo rm -Rf /var/lib/kubelet/
user@nodea:~$
```

This is a good remedy for components that will not restart due to preexisting state that is out of synch with the rest of the cluster.

Recheck that all of the Kubernetes services are stopped.

```
user@nodea:~$ pidof etcd kube-apiserver kubelet
user@nodea:~$
```

In addition to k8s components, we need to clean up what Docker is currently running.

```
user@nodea:~$ docker container rm $(docker container stop $(docker container ls -qa))
...
user@nodea:~$
```

2. Using files

When you run the `kubelet` you can supply a single manifest file (or several in a directory) as a command line argument. On `kubelet` startup these manifests will start prior to the manifests supplied by the API server. If no kube-apiserver is supplied the `kubelet` will simply run these manifests independently.

- **--pod-manifest-path=""** - Path to the pod spec file or directory of files
- **--file-check-frequency=20s** - Duration between checking config files for new data
- **--runonce[=false]** - If true, exit after spawning pods from local manifests or remote urls (can not be used with **--api-servers** and/or **--enable-server**)

We will try running the **kubelet** stand alone with a simple pod config. First create a working directory for your configuration files:

```
user@nodea:~$ cd
user@nodea:~$
```

```
user@nodea:~$ mkdir kubelet
user@nodea:~$
```

```
user@nodea:~$ cd kubelet/
user@nodea:~/kubelet$
```

Now create a simple pod to test:

```
user@nodea:~/kubelet$ vim pod.yaml
user@nodea:~/kubelet$ cat pod.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-startup
  labels:
    app: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.7.9
    ports:
    - containerPort: 80
```

```
user@nodea:~/kubelet$
```

Now start the `kubelet` with the new pod spec file supplied as a `--pod-manifest-path` parameter:

```
user@nodea:~/kubelet$ sudo ~user/k8s/_output/bin/kubelet \
--pod-manifest-path=/home/user/kubelet/pod.yaml
```

```
I0829 11:19:35.144861 42647 feature_gate.go:144] feature gates: map[]
W0829 11:19:35.145244 42647 server.go:496] No API client: no api servers specified
I0829 11:19:35.146049 42647 client.go:72] Connecting to docker on unix:///var/run/docker.sock
I0829 11:19:35.146069 42647 client.go:92] Start docker client with request timeout=2m0s
W0829 11:19:35.150012 42647 cni.go:189] Unable to update cni config: No networks found in /etc/cni/net.d
I0829 11:19:35.163655 42647 manager.go:143] cAdvisor running in container: "/user.slice"
W0829 11:19:35.184143 42647 manager.go:151] unable to connect to Rkt api service: rkt: cannot tcp Dial rkt api
service: dial tcp [::1]:15441: getsockopt: connection refused
I0829 11:19:35.193833 42647 fs.go:117] Filesystem partitions: map[/dev/sda1:{mountpoint:/var/lib/docker/aufs
major:8 minor:1 fsType:ext4 blockSize:0}]
I0829 11:19:35.197290 42647 manager.go:198] Machine: {NumCores:2 CpuFrequency:2711681 MemoryCapacity:4124880896
MachineID:6e883acc04fc7db3713776be57a3dac9 SystemUUID:F2564D56-3460-443D-57E3-836F703215A2 BootID:e21a0ef9-85e0-
4ef1-b0b9-f85c262ea596 Filesystems:[{Device:/dev/sda1 DeviceMajor:8 DeviceMinor:1 Capacity:18889830400 Type:vfs
Inodes:1179648 HasInodes:true}] DiskMap:map[8:0:{Name:sda Major:8 Minor:0 Size:21474836480 Scheduler:deadline}]
NetworkDevices:[{Name:ens33 MacAddress:00:0c:29:32:15:a2 Speed:1000 Mtu:1500}] Topology:[{Id:0 Memory:4124880896
Cores:[{Id:0 Threads:[0] Caches:[{Size:32768 Type:Data Level:1} {Size:32768 Type:Instruction Level:1} {Size:262144
Type:Unified Level:2}]}]}] Caches:[{Size:8388608 Type:Unified Level:3}] {Id:2 Memory:0 Cores:[{Id:0 Threads:[1]
Caches:[{Size:32768 Type:Data Level:1} {Size:32768 Type:Instruction Level:1} {Size:262144 Type:Unified Level:2}]}]}]
Caches:[{Size:8388608 Type:Unified Level:3}]}] CloudProvider:Unknown InstanceType:Unknown InstanceID:None}
I0829 11:19:35.198294 42647 manager.go:204] Version: {KernelVersion:4.4.0-93-generic ContainerOsVersion:Ubuntu
16.04.1 LTS DockerVersion:17.06.1-ce DockerAPIVersion:1.30 CadvisorVersion: CadvisorRevision:}
W0829 11:19:35.200357 42647 server.go:356] No api server defined - no events will be sent to API server.
I0829 11:19:35.200373 42647 server.go:536] --cgroups-per-qos enabled, but --cgroup-root was not specified.
defaulting to /
W0829 11:19:35.202322 42647 container_manager_linux.go:216] Running with swap on is not supported, please
disable swap! This will be a fatal error by default starting in K8s v1.6! In the meantime, you can opt-in to
making this a fatal error by enabling --experimental-fail-swap-on.
I0829 11:19:35.202384 42647 container_manager_linux.go:246] container manager verified user specified cgroup-
root exists: /
I0829 11:19:35.202399 42647 container_manager_linux.go:251] Creating Container Manager object based on Node
Config: {RuntimeCgroupsName: SystemCgroupsName: KubeletCgroupsName: ContainerRuntime:docker CgroupsPerQOS:true
CgroupRoot:/ CgroupDriver:cgroupfs ProtectKernelDefaults:false NodeAllocatableConfig:{KubeReservedCgroupName:
SystemReservedCgroupName: EnforceNodeAllocatable:map[pods:{}] KubeReserved:map[] SystemReserved:map[]
HardEvictionThresholds:[{Signal:memory.available Operator:LessThan Value:{Quantity:100Mi Percentage:0}}
```

```

GracePeriod:0s MinReclaim:<nil>} {Signal:nodefs.available Operator:LessThan Value:{Quantity:<nil> Percentage:0.1}
GracePeriod:0s MinReclaim:<nil>} {Signal:nodefs.inodesFree Operator:LessThan Value:{Quantity:<nil>
Percentage:0.05} GracePeriod:0s MinReclaim:<nil>}}} ExperimentalQOSReserved:map[]
I0829 11:19:35.202698 42647 kubelet.go:263] Adding manifest file: /home/user/kubelet/pod.yaml
W0829 11:19:35.206451 42647 kubelet_network.go:70] Hairpin mode set to "promiscuous-bridge" but kubenet is not
enabled, falling back to "hairpin-veth"
I0829 11:19:35.206475 42647 kubelet.go:508] Hairpin mode set to "hairpin-veth"
W0829 11:19:35.210581 42647 cni.go:189] Unable to update cni config: No networks found in /etc/cni/net.d
I0829 11:19:35.218624 42647 docker_service.go:208] Docker cri networking managed by kubernetes.io/no-op
I0829 11:19:35.231203 42647 docker_service.go:225] Setting cgroupDriver to cgroupfs
I0829 11:19:35.245809 42647 remote_runtime.go:42] Connecting to runtime service unix:///var/run/dockerhim.sock
I0829 11:19:35.247785 42647 kuberuntime_manager.go:163] Container runtime docker initialized, version: 17.06.1-
ce, apiVersion: 1.30.0
I0829 11:19:35.248966 42647 server.go:943] Started kubelet v1.7.4+793658f2d7ca7
E0829 11:19:35.249278 42647 kubelet.go:1229] Image garbage collection failed once. Stats initialization may not
have completed yet: unable to find data for container /
W0829 11:19:35.249340 42647 kubelet.go:1313] No api server defined - no node status update will be sent.
I0829 11:19:35.249581 42647 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
I0829 11:19:35.250491 42647 server.go:132] Starting to listen on 0.0.0.0:10250
I0829 11:19:35.251336 42647 server.go:310] Adding debug handlers to kubelet server.
E0829 11:19:35.258398 42647 kubelet.go:1729] Failed to check if disk space is available for the runtime: failed
to get fs info for "runtime": unable to find data for container /
E0829 11:19:35.258705 42647 kubelet.go:1737] Failed to check if disk space is available on the root partition:
failed to get fs info for "root": unable to find data for container /
E0829 11:19:35.265539 42647 docker_sandbox.go:239] Failed to stop sandbox
"b2d1c612ee8b1e63c8b3ed856c5d6e1f3af5ed9aded5b055e915b589132e8f33": Error response from daemon: {"message":"No
such container: b2d1c612ee8b1e63c8b3ed856c5d6e1f3af5ed9aded5b055e915b589132e8f33"}
I0829 11:19:35.268110 42647 fs_resource_analyzer.go:66] Starting FS ResourceAnalyzer
I0829 11:19:35.268581 42647 status_manager.go:136] Kubernetes client is nil, not starting status manager.
I0829 11:19:35.272098 42647 kubelet.go:1809] Starting kubelet main sync loop.
I0829 11:19:35.272150 42647 kubelet.go:1820] skipping pod synchronization - [container runtime is down PLEG is
not healthy: pleg was last seen active 2562047h47m16.854775807s ago; threshold is 3m0s]
W0829 11:19:35.269002 42647 container_manager_linux.go:747] CPUAccounting not enabled for pid: 42647
W0829 11:19:35.272336 42647 container_manager_linux.go:750] MemoryAccounting not enabled for pid: 42647
I0829 11:19:35.269016 42647 volume_manager.go:245] Starting Kubelet Volume Manager
E0829 11:19:35.269045 42647 container_manager_linux.go:543] [ContainerManager]: Fail to get rootfs information
unable to find data for container /
I0829 11:19:35.286710 42647 factory.go:351] Registering Docker factory
W0829 11:19:35.286739 42647 manager.go:247] Registration of the rkt container factory failed: unable to
communicate with Rkt api service: rkt: cannot tcp Dial rkt api service: dial tcp [::1]:15441: getsockopt:
connection refused
I0829 11:19:35.286802 42647 factory.go:54] Registering systemd factory

```

```

I0829 11:19:35.287007    42647 factory.go:86] Registering Raw factory
I0829 11:19:35.287741    42647 manager.go:1121] Started watching for new ooms in manager
I0829 11:19:35.288750    42647 oomparser.go:185] oomparser using systemd
I0829 11:19:35.289194    42647 manager.go:288] Starting recovery of all containers
I0829 11:19:35.360916    42647 manager.go:293] Recovery completed
I0829 11:19:35.447129    42647 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
E0829 11:19:35.450250    42647 helpers.go:771] Could not find capacity information for resource
storage.kubernetes.io/scratch
W0829 11:19:35.450297    42647 helpers.go:782] eviction manager: no observation found for eviction signal
allocatableNodeFs.available
I0829 11:19:40.279325    42647 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
I0829 11:19:40.290495    42647 kuberuntime_manager.go:457] Container {Name:nginx Image:nginx:1.7.9 Command:[] Args:
[] WorkingDir: Ports:[{Name: HostPort:0 ContainerPort:80 Protocol:TCP HostIP:}] EnvFrom:[] Env:[] Resources:
{Limits:map[] Requests:map[]} VolumeMounts:[] LivenessProbe:nil ReadinessProbe:nil Lifecycle:nil
TerminationMessagePath:/dev/termination-log TerminationMessagePolicy:File ImagePullPolicy:IfNotPresent
SecurityContext:nil Stdin:false StdinOnce:false TTY:false} is dead, but RestartPolicy says that we should restart
it.
...

```

Give Docker enough time to pull the nginx image, then list the running containers in a new shell.

```

user@nodea:~$ docker container ls --no-trunc --format "table {{.Image}}"

IMAGE
nginx@sha256:e3456c851a152494c3e4ff5fcc26f240206abac0c9d794affb40e0714846c451
gcr.io/google_containers/pause-amd64:3.0

user@nodea:~$

```

Note the nginx image tag (in the YAML), it is version 1.7.9 as requested in the spec but Kubernetes and Docker 17.06 track images by the content addressable SHA hash. Imagine we would like to change the version. Instead of manipulating the Docker containers directly, we will update the pod configuration and let the **kubelet** redeploy the new version of nginx.

In a separate shell, update your config to request image tag 1.9.1, leaving your **kubelet** running:

```

user@nodea:~$ cd kubelet/

```

```
user@nodea:~/kubelet$
```

```
user@nodea:~/kubelet$ vim pod.yaml
user@nodea:~/kubelet$ cat pod.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-startup
  labels:
    app: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.9.1
    ports:
    - containerPort: 80
```

```
user@nodea:~/kubelet$
```

Shortly after you save your file changes you should notice the following log output in the `kubelet` log:

```
...
E0829 11:29:32.360602    42647 file.go:72] unable to read config path "/home/user/kubelet/pod.yaml": error while
processing event ("/home/user/kubelet/pod.yaml": 0x400 == IN_DELETE_SELF): the watched path is deleted

W0829 11:29:35.273487    42647 container_manager_linux.go:747] CPUAccounting not enabled for pid: 42647
W0829 11:29:35.274890    42647 container_manager_linux.go:750] MemoryAccounting not enabled for pid: 42647
I0829 11:29:37.413328    42647 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
I0829 11:29:47.444459    42647 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
I0829 11:29:52.363954    42647 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
I0829 11:29:52.371401    42647 kuberuntime_manager.go:457] Container {Name:nginx Image:nginx:1.9.1 Command:[] Args:
[] WorkingDir: Ports:[{Name: HostPort:0 ContainerPort:80 Protocol:TCP HostIP:}] EnvFrom:[] Env:[] Resources:
{Limits:map[] Requests:map[]} VolumeMounts:[] LivenessProbe:nil ReadinessProbe:nil Lifecycle:nil
TerminationMessagePath:/dev/termination-log TerminationMessagePolicy:File ImagePullPolicy:IfNotPresent
SecurityContext:nil Stdin:false StdinOnce:false TTY:false} is dead, but RestartPolicy says that we should restart
```

```
it.  
E0829 11:29:53.281663    42647 kuberuntime_container.go:59] Can't make a ref to pod "nginx-startup-  
nodea_default(89beb3e314b1eb83a626cea6dcd92be7)", container nginx: selfLink was empty, can't make reference  
W0829 11:29:53.291901    42647 pod_container_deletor.go:77] Container  
"7a356a64bcc4e23b6e5e109e76b586775b871614632eef2ec0bed94270556f73" not found in pod's containers  
...
```

It may take Docker some time to pull the image. You can see the pull status by issuing the appropriate `docker pull` command:

```
user@nodea:~/kubelet$ docker image pull nginx:1.9.1  
  
1.9.1: Pulling from library/nginx  
5641bf7f839b: Pull complete  
a3ed95caeb02: Pull complete  
d003dd0d7f8a: Pull complete  
c5dd085dcc7c: Pull complete  
d95a07673dd5: Pull complete  
cec5c5855afe: Pull complete  
b315c6f2ccf3: Pull complete  
Digest: sha256:2f68b99bc0d6d25d0c56876b924ec20418544ff28e1fb89a4c27679a40da811b  
Status: Downloaded newer image for nginx:1.9.1  
  
user@nodea:~/kubelet$
```

or

```
user@nodea:~/kubelet$ docker image pull nginx:1.9.1  
  
1.9.1: Pulling from library/nginx  
5641bf7f839b: Already exists  
a3ed95caeb02: Already exists  
d003dd0d7f8a: Already exists  
c5dd085dcc7c: Already exists  
d95a07673dd5: Already exists  
cec5c5855afe: Already exists  
b315c6f2ccf3: Already exists  
Digest: sha256:2f68b99bc0d6d25d0c56876b924ec20418544ff28e1fb89a4c27679a40da811b  
Status: Image is up to date for nginx:1.9.1
```



```
user@nodea:~/kubelet$
```

or a mixture of the previous two outputs.

If you list the running containers you will see that the kubelet has started the new nginx container.

```
user@nodea:~/kubelet$ docker container ls --no-trunc --format "table {{.Image}}\t{{.CreatedAt}}\t{{.Status}}"

IMAGE                                                                                                     CREATED AT
STATUS
nginx@sha256:2f68b99bc0d6d25d0c56876b924ec20418544ff28e1fb89a4c27679a40da811b  2017-08-29 11:30:03 -0700 PDT    Up
4 minutes
gcr.io/google_containers/pause-amd64:3.0                                     2017-08-29 11:29:52 -0700 PDT    Up
4 minutes

user@nodea:~/kubelet$
```

- Note that the new 1.9 nginx container has a different SHA hash than the 1.7 version.

While the above example is just an experiment, this is exactly the way a Kubernetes master is typically boot strapped. For example, the kubeadm installer, configures a kubelet with a pod manifest path and then adds pod specs to the directory for etcd, the api-server, the controller-manager, and the scheduler. In this way the kubelet is the only Kubernetes service actually running on the host, all of the other services run in containers.

The kubelet in turn is generally configured as a systemd service, so if the kubelet fails, systemd will restart it.

3. HTTP endpoint

The `kubelet` also has the ability load manifests from URL based resources. The following switches configure this feature:

- `--manifest-url=""` - URL for accessing the container manifest
- `--manifest-url-header=""` - HTTP header to use when accessing the manifest URL, with the key separated from the value with a ':', as in 'key:value'
- `--http-check-frequency=20s` - Duration between checking http for new data

Let's try running the `kubelet` using a config supplied by URL.

Stop the `kubelet` with ^C.

Clean up the running containers.

```
user@nodea:~/kubelet$ docker container rm $(docker container stop $(docker container ls -qa))  
...  
user@nodea:~/kubelet$
```

Clear all of the kubelet state:

```
user@nodea:~/kubelet$ sudo rm -Rf /var/lib/kubelet  
user@nodea:~/kubelet$
```

The Kubernetes GitHub repo has a sample pod we can use for this test:

```
user@nodea:~/kubelet$ curl https://raw.githubusercontent.com/kubernetes/kubernetes/master/examples/pod  
  
# Copy of pod.yaml without file extension for test  
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx  
  labels:  
    name: nginx  
spec:  
  containers:  
  - name: nginx  
    image: nginx  
    ports:  
    - containerPort: 80  
  
user@nodea:~/kubelet$
```

Now rerun the `kubelet`, providing it the URL from the Kubernetes repo.

```
user@nodea:~/kubelet$ sudo ~user/k8s/_output/bin/kubelet \  
--manifest-url=https://raw.githubusercontent.com/kubernetes/kubernetes/master/examples/pod  
  
I0829 11:40:29.274971    44070 feature_gate.go:144] feature gates: map[]
```

```

W0829 11:40:29.275107 44070 server.go:496] No API client: no api servers specified
I0829 11:40:29.275152 44070 client.go:72] Connecting to docker on unix:///var/run/docker.sock
I0829 11:40:29.275161 44070 client.go:92] Start docker client with request timeout=2m0s
W0829 11:40:29.277792 44070 cni.go:189] Unable to update cni config: No networks found in /etc/cni/net.d
I0829 11:40:29.283748 44070 manager.go:143] cAdvisor running in container: "/user.slice"
W0829 11:40:29.296253 44070 manager.go:151] unable to connect to Rkt api service: rkt: cannot tcp Dial rkt api
service: dial tcp [::1]:15441: getsockopt: connection refused
I0829 11:40:29.308595 44070 fs.go:117] Filesystem partitions: map[/dev/sda1:{mountpoint:/var/lib/docker/aufs
major:8 minor:1 fsType:ext4 blockSize:0}]
I0829 11:40:29.309828 44070 manager.go:198] Machine: {NumCores:2 CpuFrequency:2711681 MemoryCapacity:4124880896
MachineID:6e883acc04fc7db3713776be57a3dac9 SystemUUID:F2564D56-3460-443D-57E3-836F703215A2 BootID:e21a0ef9-85e0-
4ef1-b0b9-f85c262ea596 Filesystems:[{Device:/dev/sda1 DeviceMajor:8 DeviceMinor:1 Capacity:18889830400 Type:vfs
Inodes:1179648 HasInodes:true}] DiskMap:map[8:0:{Name:sda Major:8 Minor:0 Size:21474836480 Scheduler:deadline}]
NetworkDevices:[{Name:ens33 MacAddress:00:0c:29:32:15:a2 Speed:1000 Mtu:1500}] Topology:[{Id:0 Memory:4124880896
Cores:[{Id:0 Threads:[0] Caches:[{Size:32768 Type:Data Level:1} {Size:32768 Type:Instruction Level:1} {Size:262144
Type:Unified Level:2}]}]} Caches:[{Size:8388608 Type:Unified Level:3}]} {Id:2 Memory:0 Cores:[{Id:0 Threads:[1]
Caches:[{Size:32768 Type:Data Level:1} {Size:32768 Type:Instruction Level:1} {Size:262144 Type:Unified Level:2}]}]}
Caches:[{Size:8388608 Type:Unified Level:3}]}] CloudProvider:Unknown InstanceType:Unknown InstanceID:None}
I0829 11:40:29.310868 44070 manager.go:204] Version: {KernelVersion:4.4.0-93-generic ContainerOsVersion:Ubuntu
16.04.1 LTS DockerVersion:17.06.1-ce DockerAPIVersion:1.30 CadvisorVersion:CadvisorRevision:}
W0829 11:40:29.311508 44070 server.go:356] No api server defined - no events will be sent to API server.
I0829 11:40:29.312010 44070 server.go:536] --cgroups-per-qos enabled, but --cgroup-root was not specified.
defaulting to /
W0829 11:40:29.313661 44070 container_manager_linux.go:216] Running with swap on is not supported, please
disable swap! This will be a fatal error by default starting in K8s v1.6! In the meantime, you can opt-in to
making this a fatal error by enabling --experimental-fail-swap-on.
I0829 11:40:29.313724 44070 container_manager_linux.go:246] container manager verified user specified cgroup-
root exists: /
I0829 11:40:29.313744 44070 container_manager_linux.go:251] Creating Container Manager object based on Node
Config: {RuntimeCgroupsName: SystemCgroupsName: KubeletCgroupsName: ContainerRuntime:docker CgroupsPerQOS:true
CgroupRoot:/ CgroupDriver:cgroupfs ProtectKernelDefaults:false NodeAllocatableConfig:{KubeReservedCgroupName:
SystemReservedCgroupName: EnforceNodeAllocatable:map[pods:{}] KubeReserved:map[] SystemReserved:map[]
HardEvictionThresholds:[{Signal:memory.available Operator:LessThan Value:{Quantity:100Mi Percentage:0}
GracePeriod:0s MinReclaim:<nil>} {Signal:nodfs.available Operator:LessThan Value:{Quantity:<nil> Percentage:0.1}
GracePeriod:0s MinReclaim:<nil>} {Signal:nodfs.inodesFree Operator:LessThan Value:{Quantity:<nil>
Percentage:0.05} GracePeriod:0s MinReclaim:<nil>}}} ExperimentalQOSReserved:map[]}
I0829 11:40:29.314036 44070 kubelet.go:269] Adding manifest url
"https://raw.githubusercontent.com/kubernetes/kubernetes/master/examples/pod" with HTTP header map[]
W0829 11:40:29.316698 44070 kubelet_network.go:70] Hairpin mode set to "promiscuous-bridge" but kubenet is not
enabled, falling back to "hairpin-veth"
I0829 11:40:29.316792 44070 kubelet.go:508] Hairpin mode set to "hairpin-veth"
W0829 11:40:29.319984 44070 cni.go:189] Unable to update cni config: No networks found in /etc/cni/net.d
I0829 11:40:29.328644 44070 docker_service.go:208] Docker cri networking managed by kubernetes.io/no-op

```

```

I0829 11:40:29.341116 44070 docker_service.go:225] Setting cgroupDriver to cgroupfs
I0829 11:40:29.355454 44070 remote_runtime.go:42] Connecting to runtime service unix:///var/run/dockershim.sock
I0829 11:40:29.357024 44070 kubernetemanager.go:163] Container runtime docker initialized, version: 17.06.1-
ce, apiVersion: 1.30.0
I0829 11:40:29.358345 44070 server.go:943] Started kubelet v1.7.4+793658f2d7ca7
E0829 11:40:29.358549 44070 kubelet.go:1229] Image garbage collection failed once. Stats initialization may not
have completed yet: unable to find data for container /
W0829 11:40:29.358620 44070 kubelet.go:1313] No api server defined - no node status update will be sent.
I0829 11:40:29.358824 44070 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
I0829 11:40:29.359053 44070 server.go:132] Starting to listen on 0.0.0.0:10250
I0829 11:40:29.359851 44070 server.go:310] Adding debug handlers to kubelet server.
E0829 11:40:29.364089 44070 kubelet.go:1729] Failed to check if disk space is available for the runtime: failed
to get fs info for "runtime": unable to find data for container /
E0829 11:40:29.366864 44070 kubelet.go:1737] Failed to check if disk space is available on the root partition:
failed to get fs info for "root": unable to find data for container /
I0829 11:40:29.367563 44070 fs_resource_analyzer.go:66] Starting FS ResourceAnalyzer
I0829 11:40:29.367595 44070 status_manager.go:136] Kubernetes client is nil, not starting status manager.
I0829 11:40:29.367604 44070 kubelet.go:1809] Starting kubelet main sync loop.
I0829 11:40:29.367627 44070 kubelet.go:1820] skipping pod synchronization - [container runtime is down PLEG is
not healthy: pleg was last seen active 2562047h47m16.854775807s ago; threshold is 3m0s]
W0829 11:40:29.368160 44070 container_manager_linux.go:747] CPUAccounting not enabled for pid: 44070
W0829 11:40:29.368750 44070 container_manager_linux.go:750] MemoryAccounting not enabled for pid: 44070
E0829 11:40:29.368357 44070 container_manager_linux.go:543] [ContainerManager]: Fail to get rootfs information
unable to find data for container /
I0829 11:40:29.368386 44070 volume_manager.go:245] Starting Kubelet Volume Manager
E0829 11:40:29.376221 44070 docker_sandbox.go:239] Failed to stop sandbox
"41ac78587353d8123c00e20352c4191e9c609e2fc3306142b5c0da07e9076f13": Error response from daemon: {"message":"No
such container: 41ac78587353d8123c00e20352c4191e9c609e2fc3306142b5c0da07e9076f13"}
I0829 11:40:29.390103 44070 factory.go:351] Registering Docker factory
W0829 11:40:29.390359 44070 manager.go:247] Registration of the rkt container factory failed: unable to
communicate with Rkt api service: rkt: cannot tcp Dial rkt api service: dial tcp [::1]:15441: getsockopt:
connection refused
I0829 11:40:29.390541 44070 factory.go:54] Registering systemd factory
I0829 11:40:29.390962 44070 factory.go:86] Registering Raw factory
I0829 11:40:29.391326 44070 manager.go:1121] Started watching for new ooms in manager
I0829 11:40:29.391732 44070 oomparser.go:185] oomparser using systemd
I0829 11:40:29.392138 44070 manager.go:288] Starting recovery of all containers
I0829 11:40:29.465837 44070 manager.go:293] Recovery completed
I0829 11:40:29.558781 44070 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
E0829 11:40:29.561975 44070 helpers.go:771] Could not find capacity information for resource
storage.kubernetes.io/scratch

```

```

W0829 11:40:29.562025    44070 helpers.go:782] eviction manager: no observation found for eviction signal
allocatableNodeFs.available
I0829 11:40:34.369250    44070 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
W0829 11:40:34.371578    44070 pod_container_deletor.go:77] Container
"41ac78587353d8123c00e20352c4191e9c609e2fc3306142b5c0da07e9076f13" not found in pod's containers
I0829 11:40:34.377291    44070 kuberuntime_manager.go:457] Container {Name:nginx Image:nginx Command:[] Args:[]
WorkingDir: Ports:[{Name: HostPort:0 ContainerPort:80 Protocol:TCP HostIP:}] EnvFrom:[] Env:[] Resources:
{Limits:map[] Requests:map[]} VolumeMounts:[] LivenessProbe:nil ReadinessProbe:nil Lifecycle:nil
TerminationMessagePath:/dev/termination-log TerminationMessagePolicy:File ImagePullPolicy:Always
SecurityContext:nil Stdin:false StdinOnce:false TTY:false} is dead, but RestartPolicy says that we should restart
it.
I0829 11:40:39.592182    44070 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
...

```

Notice the informational log output indicating that the `kubelet` is adding the desired URL to its manifest list.

```

I0829 11:40:29.314036 44070 kubelet.go:269] Adding manifest url "https://raw.githubusercontent.com/kubernetes/kubernetes/master/examples/pod"
with HTTP header map[]

```

Examine the containers running through Docker.

```

user@nodea:~/kubelet$ docker container ls --no-trunc --format "table {{.Image}}\t{{.CreatedAt}}\t{{.Status}}"

```

IMAGE	CREATED AT	
STATUS		
nginx@sha256:788fa27763db6d69ad3444e8ba72f947df9e7e163bad7c1f5614f8fd27a311c3	2017-08-29 11:40:35 -0700 PDT	Up
2 minutes		
gcr.io/google_containers/pause-amd64:3.0	2017-08-29 11:40:34 -0700 PDT	Up
2 minutes		

```

user@nodea:~/kubelet$

```

Stop the `kubelet` again with `^C`, then list the running containers.

```

user@nodea:~/kubelet$ docker container ls --no-trunc --format "table {{.Image}}\t{{.CreatedAt}}\t{{.Status}}"

```

IMAGE STATUS	CREATED AT	
nginx@sha256:788fa27763db6d69ad3444e8ba72f947df9e7e163bad7c1f5614f8fd27a311c3 3 minutes	2017-08-29 11:40:35 -0700 PDT	Up
gcr.io/google_containers/pause-amd64:3.0 3 minutes	2017-08-29 11:40:34 -0700 PDT	Up

user@nodea:~/kubelet\$

As you can see the nginx pod started by the `kubelet` is still running.

Display the labels associated with the nginx container:

```
user@nodea:~/kubelet$ docker container inspect \
$(docker container ls --filter=ancestor=nginx -q) | jq -r '.[].Config.Labels'

{
  "annotation.io.kubernetes.container.hash": "30f2a656",
  "annotation.io.kubernetes.container.ports": "[{\"containerPort\":80,\"protocol\":\"TCP\"}]",
  "annotation.io.kubernetes.container.restartCount": "0",
  "annotation.io.kubernetes.container.terminationMessagePath": "/dev/termination-log",
  "annotation.io.kubernetes.container.terminationMessagePolicy": "File",
  "annotation.io.kubernetes.pod.terminationGracePeriod": "30",
  "io.kubernetes.container.logpath": "/var/log/pods/ab4f45926942575bebaa13c947218fce/nginx_0.log",
  "io.kubernetes.container.name": "nginx",
  "io.kubernetes.docker.type": "container",
  "io.kubernetes.pod.name": "nginx-nodea",
  "io.kubernetes.pod.namespace": "default",
  "io.kubernetes.pod.uid": "ab4f45926942575bebaa13c947218fce",
  "io.kubernetes.sandbox.id": "ce2d0a2474457925ae2d4d28b1fc273954bc19faa59ab776e34365877a2a2e17"
}
user@nodea:~/kubelet$
```

As you can see, it is easy for the Kubelet to identify its own containers.

Restart the `kubelet` with no arguments:

```
user@nodea:~/kubelet$ sudo ~user/k8s/_output/bin/kubelet
```

```

I0829 11:44:57.624360 44351 feature_gate.go:144] feature gates: map[]
W0829 11:44:57.624721 44351 server.go:496] No API client: no api servers specified
I0829 11:44:57.625019 44351 client.go:72] Connecting to docker on unix:///var/run/docker.sock
I0829 11:44:57.625280 44351 client.go:92] Start docker client with request timeout=2m0s
W0829 11:44:57.626978 44351 cni.go:189] Unable to update cni config: No networks found in /etc/cni/net.d
I0829 11:44:57.641173 44351 manager.go:143] cAdvisor running in container: "/user.slice"
W0829 11:44:57.655334 44351 manager.go:151] unable to connect to Rkt api service: rkt: cannot tcp Dial rkt api
service: dial tcp [::1]:15441: getsockopt: connection refused
I0829 11:44:57.669240 44351 fs.go:117] Filesystem partitions: map[/dev/sda1:{mountpoint:/var/lib/docker/aufs
major:8 minor:1 fsType:ext4 blockSize:0}]
I0829 11:44:57.672457 44351 manager.go:198] Machine: {NumCores:2 CpuFrequency:2711681 MemoryCapacity:4124880896
MachineID:6e883acc04fc7db3713776be57a3dac9 SystemUUID:F2564D56-3460-443D-57E3-836F703215A2 BootID:e21a0ef9-85e0-
4ef1-b0b9-f85c262ea596 Filesystems:[{Device:/dev/sda1 DeviceMajor:8 DeviceMinor:1 Capacity:18889830400 Type:vfs
Inodes:1179648 HasInodes:true}] DiskMap:map[8:0:{Name:sda Major:8 Minor:0 Size:21474836480 Scheduler:deadline}]
NetworkDevices:[{Name:ens33 MacAddress:00:0c:29:32:15:a2 Speed:1000 Mtu:1500}] Topology:[{Id:0 Memory:4124880896
Cores:[{Id:0 Threads:[0] Caches:[{Size:32768 Type:Data Level:1} {Size:32768 Type:Instruction Level:1} {Size:262144
Type:Unified Level:2}]}]}] Caches:[{Size:8388608 Type:Unified Level:3}] {Id:2 Memory:0 Cores:[{Id:0 Threads:[1]
Caches:[{Size:32768 Type:Data Level:1} {Size:32768 Type:Instruction Level:1} {Size:262144 Type:Unified Level:2}]}]}]
Caches:[{Size:8388608 Type:Unified Level:3}]}] CloudProvider:Unknown InstanceType:Unknown InstanceID:None}
I0829 11:44:57.673578 44351 manager.go:204] Version: {KernelVersion:4.4.0-93-generic ContainerOsVersion:Ubuntu
16.04.1 LTS DockerVersion:17.06.1-ce DockerAPIVersion:1.30 CadvisorVersion:CadvisorRevision:}
W0829 11:44:57.674146 44351 server.go:356] No api server defined - no events will be sent to API server.
I0829 11:44:57.674176 44351 server.go:536] --cgroups-per-qos enabled, but --cgroup-root was not specified.
defaulting to /
W0829 11:44:57.675237 44351 container_manager_linux.go:216] Running with swap on is not supported, please
disable swap! This will be a fatal error by default starting in K8s v1.6! In the meantime, you can opt-in to
making this a fatal error by enabling --experimental-fail-swap-on.
I0829 11:44:57.675396 44351 container_manager_linux.go:246] container manager verified user specified cgroup-
root exists: /
I0829 11:44:57.675491 44351 container_manager_linux.go:251] Creating Container Manager object based on Node
Config: {RuntimeCgroupsName: SystemCgroupsName: KubeletCgroupsName: ContainerRuntime:docker CgroupsPerQOS:true
CgroupRoot:/ CgroupDriver:cgroupfs ProtectKernelDefaults:false NodeAllocatableConfig:{KubeReservedCgroupName:
SystemReservedCgroupName: EnforceNodeAllocatable:map[pods:{}]} KubeReserved:map[] SystemReserved:map[]
HardEvictionThresholds:[{Signal:memory.available Operator:LessThan Value:{Quantity:100Mi Percentage:0}
GracePeriod:0s MinReclaim:<nil>} {Signal:nodfs.available Operator:LessThan Value:{Quantity:<nil> Percentage:0.1}
GracePeriod:0s MinReclaim:<nil>} {Signal:nodfs.inodesFree Operator:LessThan Value:{Quantity:<nil>
Percentage:0.05} GracePeriod:0s MinReclaim:<nil>}}] ExperimentalQOSReserved:map[]}
W0829 11:44:57.677961 44351 kubelet_network.go:70] Hairpin mode set to "promiscuous-bridge" but kubenet is not
enabled, falling back to "hairpin-veth"
I0829 11:44:57.678012 44351 kubelet.go:508] Hairpin mode set to "hairpin-veth"
W0829 11:44:57.680324 44351 cni.go:189] Unable to update cni config: No networks found in /etc/cni/net.d
I0829 11:44:57.689597 44351 docker_service.go:208] Docker cri networking managed by kubernetes.io/no-op
I0829 11:44:57.707612 44351 docker_service.go:225] Setting cgroupDriver to cgroupfs

```

```

I0829 11:44:57.727897 44351 remote_runtime.go:42] Connecting to runtime service unix:///var/run/dockerhim.sock
I0829 11:44:57.729464 44351 kuberuntime_manager.go:163] Container runtime docker initialized, version: 17.06.1-
ce, apiVersion: 1.30.0
I0829 11:44:57.730939 44351 server.go:943] Started kubelet v1.7.4+793658f2d7ca7
E0829 11:44:57.730957 44351 kubelet.go:1229] Image garbage collection failed once. Stats initialization may not
have completed yet: unable to find data for container /
W0829 11:44:57.731025 44351 kubelet.go:1313] No api server defined - no node status update will be sent.
I0829 11:44:57.731102 44351 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
I0829 11:44:57.731239 44351 server.go:132] Starting to listen on 0.0.0.0:10250
I0829 11:44:57.731827 44351 server.go:310] Adding debug handlers to kubelet server.
E0829 11:44:57.735584 44351 kubelet.go:1729] Failed to check if disk space is available for the runtime: failed
to get fs info for "runtime": unable to find data for container /
E0829 11:44:57.735614 44351 kubelet.go:1737] Failed to check if disk space is available on the root partition:
failed to get fs info for "root": unable to find data for container /
I0829 11:44:57.736439 44351 fs_resource_analyzer.go:66] Starting FS ResourceAnalyzer
I0829 11:44:57.736470 44351 status_manager.go:136] Kubernetes client is nil, not starting status manager.
I0829 11:44:57.736477 44351 kubelet.go:1809] Starting kubelet main sync loop.
I0829 11:44:57.736505 44351 kubelet.go:1820] skipping pod synchronization - [container runtime is down PLEG is
not healthy: pleg was last seen active 2562047h47m16.854775807s ago; threshold is 3m0s]
I0829 11:44:57.736708 44351 volume_manager.go:245] Starting Kubelet Volume Manager
W0829 11:44:57.737327 44351 container_manager_linux.go:747] CPUAccounting not enabled for pid: 44351
W0829 11:44:57.737509 44351 container_manager_linux.go:750] MemoryAccounting not enabled for pid: 44351
E0829 11:44:57.737776 44351 container_manager_linux.go:543] [ContainerManager]: Fail to get rootfs information
unable to find data for container /
I0829 11:44:57.756019 44351 factory.go:351] Registering Docker factory
W0829 11:44:57.756298 44351 manager.go:247] Registration of the rkt container factory failed: unable to
communicate with Rkt api service: rkt: cannot tcp Dial rkt api service: dial tcp [::1]:15441: getsockopt:
connection refused
I0829 11:44:57.756487 44351 factory.go:54] Registering systemd factory
I0829 11:44:57.756996 44351 factory.go:86] Registering Raw factory
I0829 11:44:57.757476 44351 manager.go:1121] Started watching for new ooms in manager
I0829 11:44:57.759187 44351 oomparser.go:185] oomparser using systemd
I0829 11:44:57.759771 44351 manager.go:288] Starting recovery of all containers
I0829 11:44:57.821291 44351 manager.go:293] Recovery completed
I0829 11:44:57.918714 44351 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
E0829 11:44:57.921871 44351 helpers.go:771] Could not find capacity information for resource
storage.kubernetes.io/scratch
W0829 11:44:57.921945 44351 helpers.go:782] eviction manager: no observation found for eviction signal
allocatableNodeFs.available
E0829 11:45:02.852311 44351 kuberuntime_container.go:59] Can't make a ref to pod "nginx-
nodea_default(ab4f45926942575bebaa13c947218fce)", container nginx: selfLink was empty, can't make reference

```



```
W0829 11:45:02.952758    44351 docker_sandbox.go:342] failed to read pod IP from plugin/docker: Couldn't find
network status for default/nginx-nodea through plugin: invalid network status for
E0829 11:45:03.742881    44351 kuberuntime_container.go:59] Can't make a ref to pod "nginx-
nodea_default(ab4f45926942575bebaa13c947218fce)", container nginx: selfLink was empty, can't make reference
W0829 11:45:03.809963    44351 pod_container_deletor.go:77] Container
"ce2d0a2474457925ae2d4d28b1fc273954bc19faa59ab776e34365877a2a2e17" not found in pod's containers
I0829 11:45:07.948608    44351 kubelet_node_status.go:247] Setting node annotation to enable volume controller
attach/detach
...
```

Notice the Warning log output

```
"nginx-nodea_default(ab4f45926942575bebaa13c947218fce)", container nginx: selfLink was empty, can't make reference
```

at the bottom of the display. The `kubelet` is discovering containers running that it has no manifests for.

In another terminal display the running containers:

```
user@nodea:~$ docker container ls
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
user@nodea:~\$						

The `kubelet` takes ownership of the node, seriously, any pods running that the `kubelet` can not reconcile with the manifests it has been assigned are stopped and removed. If you run ad hoc containers using docker commands they will not have the Kubelet specific labels and the Kubelet will ignore them.

4. HTTP server

The `kubelet` has its own REST API and can be run as a standalone server when appropriate. You can request information including pod details (`/pods`) and overall node status (`/healthz`).

The REST endpoint on the `kubelet` is enabled by default but you can disable it with the `--enable-server=false` switch.

- `--enable-server=[true]` - Enable the kubelet's server

Try curling a list of pods from the Kubelet.

```
user@nodea:~/kubelet$ curl -s --insecure https://localhost:10250/pods | jq .
```

```
{
  "kind": "PodList",
  "apiVersion": "v1",
  "metadata": {},
  "items": null
}
user@nodea:~/kubelet$
```

Stop & rerun the `kubelet` with the previous IRI based PodSpec.

```
user@nodea:~/kubelet$ sudo ~user/k8s/_output/bin/kubelet \
--manifest-url=https://raw.githubusercontent.com/kubernetes/kubernetes/master/examples/pod
...
```

Try the pod listing again.

```
user@nodea:~/kubelet$ curl -s --insecure https://localhost:10250/pods | jq .items[].spec
```

```
{
  "containers": [
    {
      "name": "nginx",
      "image": "nginx",
      "ports": [
        {
          "containerPort": 80,
          "protocol": "TCP"
        }
      ],
      "resources": {},
      "terminationMessagePath": "/dev/termination-log",
      "terminationMessagePolicy": "File",
      "imagePullPolicy": "Always"
    }
  ],
  "restartPolicy": "Always",
  "terminationGracePeriodSeconds": 30,
  "dnsPolicy": "ClusterFirst",

```

```
"nodeName": "nodea",
"securityContext": {},
"schedulerName": "default-scheduler"
}
```

```
user@nodea:~/kubelet$
```

Next stop the kubelet and rerun it disabling the HTTP server with `--enable-server=false`.

```
user@nodea:~$ sudo ~user/k8s/_output/bin/kubelet \
--manifest-url=https://raw.githubusercontent.com/kubernetes/kubernetes/master/examples/pod \
--enable-server=false

...
```

The `kubelet` is running and our pod is started but the REST endpoint is down.

```
user@nodea:~/kubelet$ curl -svk https://localhost:10250/pods

* Trying ::1...
* connect to ::1 port 10250 failed: Connection refused
* Trying 127.0.0.1...
* connect to 127.0.0.1 port 10250 failed: Connection refused
* Failed to connect to localhost port 10250: Connection refused
* Closing connection 0

user@nodea:~/kubelet$
```

```
user@nodea:~/kubelet$ docker container ls
```

CONTAINER ID	IMAGE	COMMAND	CREATED
7a0b1acd0a95	nginx	"nginx -g 'daemon ...'"	About a minute ago Up
About a minute	k8s_nginx_nginx-nodea_default_ab4f45926942575bebaa13c947218fce_0		
17fdc80a4c01	gcr.io/google_containers/pause-amd64:3.0	"/pause"	About a minute ago Up
About a minute	k8s_POD_nginx-nodea_default_ab4f45926942575bebaa13c947218fce_0		

```
user@nodea:~/kubelet$
```

Restart the `kubelet` with HTTP enabled (remove the `--enable-server` or set it to `true`)

```
user@nodea:~/kubelet$ sudo ~user/k8s/_output/bin/kubelet \
--manifest-url=https://raw.githubusercontent.com/kubernetes/kubernetes/master/examples/pod \
--enable-server=true
...
```

Try to stop the nginx container owned by the kubelet via Docker.

```
user@nodea:~/kubelet$ docker container ls
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
7a0b1acd0a95	nginx	"nginx -g 'daemon ...'"	2 minutes ago	Up 2
17fdc80a4c01	gcr.io/google_containers/pause-amd64:3.0	"/pause"	2 minutes ago	Up 2

```
user@nodea:~/kubelet$
```

```
user@nodea:~/kubelet$ docker container kill $(docker container ls --filter=ancestor=nginx -q)
```

```
7a0b1acd0a95
```

```
user@nodea:~/kubelet$
```

```
user@nodea:~/kubelet$ docker container ls
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
da140e7f5d1b	nginx	"nginx -g 'daemon ...'"	3 seconds ago	Up 2

```
17fdc80a4c01      gcr.io/google_containers/pause-amd64:3.0    "/pause"      2 minutes ago    Up 2
minutes          k8s_POD_nginx-nodea_default_ab4f45926942575bebaa13c947218fce_0

user@nodea:~/kubelet$
```

- What happened?

Docker reports that it killed the container in question. However a new `docker container ls` shows the same nginx image running. However, if you look carefully, you will see that it is *not* the same container. You killed one container (7a0b1acd0a95 in the example) and the `kubelet` started a new copy of the image (container da140e7f5d1b in the example). The `kubelet` will *never* restart a container, it will only run new copies of the image when an old container fails.

Look at the `kubelet` log output for clues.

When the container fails, the `kubelet` checks the backoff time and if it has expired the `kubelet` tries to recreate the container. The back off ensures that the `kubelet` will not try to restart the container more than once in the backoff time window.

This behavior is consistent with the general Kubernetes philosophy, users supply the desired state and Kubernetes ensures that it is enforced as the actual state. As long as this `kubelet` has the podspec for nginx, it will make sure nginx is running.

5. Health check

The `kubelet` offers a basic health check endpoint which is used to verify reachability and liveness of the `kubelet`.

The `/healthz` path can be curled easily, try it:

```
user@nodea:~/kubelet$ curl -v 127.0.0.1:10248/healthz && echo

* Trying 127.0.0.1...
* Connected to 127.0.0.1 (127.0.0.1) port 10248 (#0)
> GET /healthz HTTP/1.1
> Host: 127.0.0.1:10248
> User-Agent: curl/7.47.0
> Accept: */*
>
< HTTP/1.1 200 OK
< Date: Tue, 29 Aug 2017 18:51:34 GMT
< Content-Length: 2
< Content-Type: text/plain; charset=utf-8
<
```

```
* Connection #0 to host 127.0.0.1 left intact
ok
```

```
user@nodea:~/kubelet$
```

Be advised that this is a very primitive health check, it only tells you that the `kubelet` is running. You can stop the Docker daemon (crash all pods) and the `kubelet` will still return ok. This only tells you that the kubelet is ok, it says nothing about the rest of the node.

6. Spec

You can use the spec endpoint to retrieve general information about this kubelet's node.

Try it:

```
user@nodea:~/kubelet$ curl -sL 127.0.0.1:10255/spec | jq .
```

```
{
  "num_cores": 2,
  "cpu_frequency_khz": 2711681,
  "memory_capacity": 4124880896,
  "machine_id": "6e883acc04fc7db3713776be57a3dac9",
  "system_uuid": "F2564D56-3460-443D-57E3-836F703215A2",
  "boot_id": "e21a0ef9-85e0-4ef1-b0b9-f85c262ea596",
  "filesystems": [
    {
      "device": "/dev/sda1",
      "capacity": 18889830400,
      "type": "vfs",
      "inodes": 1179648,
      "has_inodes": true
    }
  ],
  "disk_map": {
    "8:0": {
      "name": "sda",
      "major": 8,
      "minor": 0,
      "size": 21474836480,
      "scheduler": "deadline"
    }
  }
},
```

```
"network_devices": [
  {
    "name": "ens33",
    "mac_address": "00:0c:29:32:15:a2",
    "speed": 1000,
    "mtu": 1500
  }
],
"topology": [
  {
    "node_id": 0,
    "memory": 4124880896,
    "cores": [
      {
        "core_id": 0,
        "thread_ids": [
          0
        ],
        "caches": [
          {
            "size": 32768,
            "type": "Data",
            "level": 1
          },
          {
            "size": 32768,
            "type": "Instruction",
            "level": 1
          },
          {
            "size": 262144,
            "type": "Unified",
            "level": 2
          }
        ]
      }
    ]
  }
],
"caches": [
  {
    "size": 8388608,
    "type": "Unified",
    "level": 3
  }
]
```

```

    ]
  },
  {
    "node_id": 2,
    "memory": 0,
    "cores": [
      {
        "core_id": 0,
        "thread_ids": [
          1
        ],
        "caches": [
          {
            "size": 32768,
            "type": "Data",
            "level": 1
          },
          {
            "size": 32768,
            "type": "Instruction",
            "level": 1
          },
          {
            "size": 262144,
            "type": "Unified",
            "level": 2
          }
        ]
      }
    ],
    "caches": [
      {
        "size": 8388608,
        "type": "Unified",
        "level": 3
      }
    ]
  }
],
"cloud_provider": "Unknown",
"instance_type": "Unknown",
"instance_id": "None"
}

```



```
user@nodea:~/kubelet$
```

Congratulations you have successfully completed the kubelet lab!

Copyright (c) 2013-2017 RX-M LLC, Cloud Native Consulting, all rights reserved