

Kubernetes

Foundation

RX-M Cloud Native Advisory, Consulting and Training

2

Copyright 2013-2017, RX-M LLC

▪ Microservice Oriented

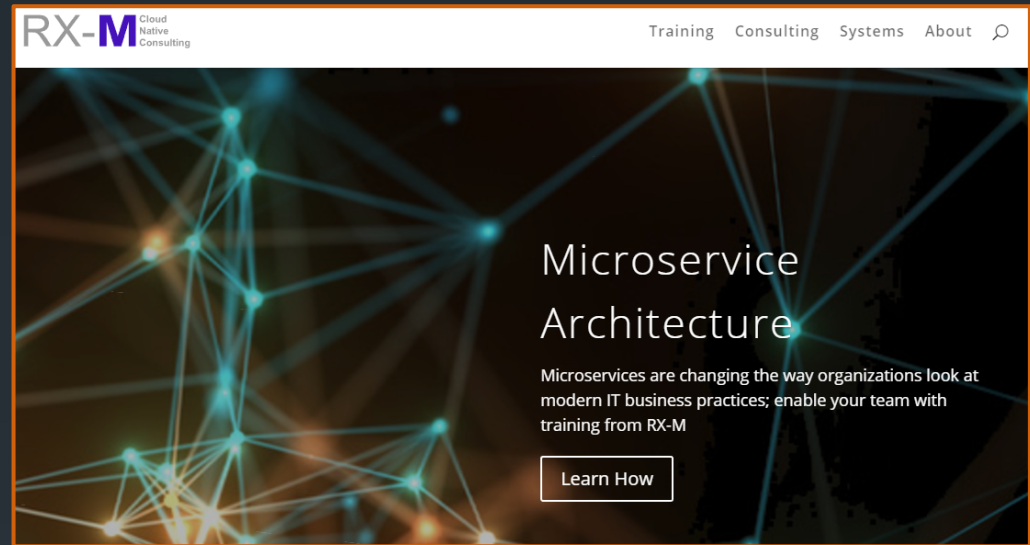
- Microservices Foundation [3 Day]
- Building Microservices on AWS [3 Day]
- Building Microservices with Go [3 Day]
- Building Microservices with Thrift [2 Day]
- Building Microservices with gRPC [2 Day]

▪ Container Packaged

- Docker Foundation [3 Day]
- Docker Advanced [2 Day]
- OCI [2 Day]
- CNI [2 Day]
- Containerd [2 Day]
- Rocket [2 Day]

▪ Dynamically Managed

- Cloud Native Container Networking and Cisco ACI [3 Day]
- Docker Orchestration (Compose/Swarm) [2 Day]
- Kubernetes Foundation [2 Day]
- Kubernetes Advanced [3 Day]
- Mesos Foundation [2 Day]
- MANTL [2 Day]
- Nomad [2 Day]



RX-M Cloud
Native
Consulting

5: Controllers

Objectives

- Understanding labels and selectors and how they are used to identify resources in the cluster
- Explore the types, roles and workings of Controllers
 - Deployments, Replica Sets, Replication Controllers, DaemonSets, Jobs, CronJobs, StatefulSets, PetSets, HorizontalPodAutoscaler
 - Understand the differences between RSs & RCs
- Explore additional spec config options
 - Policies
 - Health and readiness
 - Pre and post actions

Controller Types

5

Copyright 2013-2017, RX-M LLC

- While pods can be created directly, in a cluster they are almost always created through a Controller
- **Deployments** – for pods that back long running services with support for upgrades and scaling
 - Deployment object describes a desired state
 - Deployment controller changes the actual state to the desired state at a controlled rate
- **Replication Controller** (RC) – ensures that a specified number of Pod replicas are running at any one time
- **Replica Set** (RS) – the next-generation Replication Controller
 - Deployments recommended instead of directly using ReplicaSets, unless you require custom update orchestration or don't require updates
 - You may never need to manipulate ReplicaSet objects
 - Only difference between a RS and a RC right now is the selector support
- **Daemon Set** – ensures that all (or some) Nodes run a copy of a Pod
- **Job** – for one shot pods used to complete a start to finish task
- **CronJob** – manages time-based Jobs
- **StatefulSets** – for pods with an identity tied to state storage in volumes or other backends
 - **PetSets** – renamed StatefulSet starting in Kubernetes version 1.5
- **HorizontalPodAutoscaler** – automatically scales the number of pods in a RC, Deployment or RS based on observed metrics



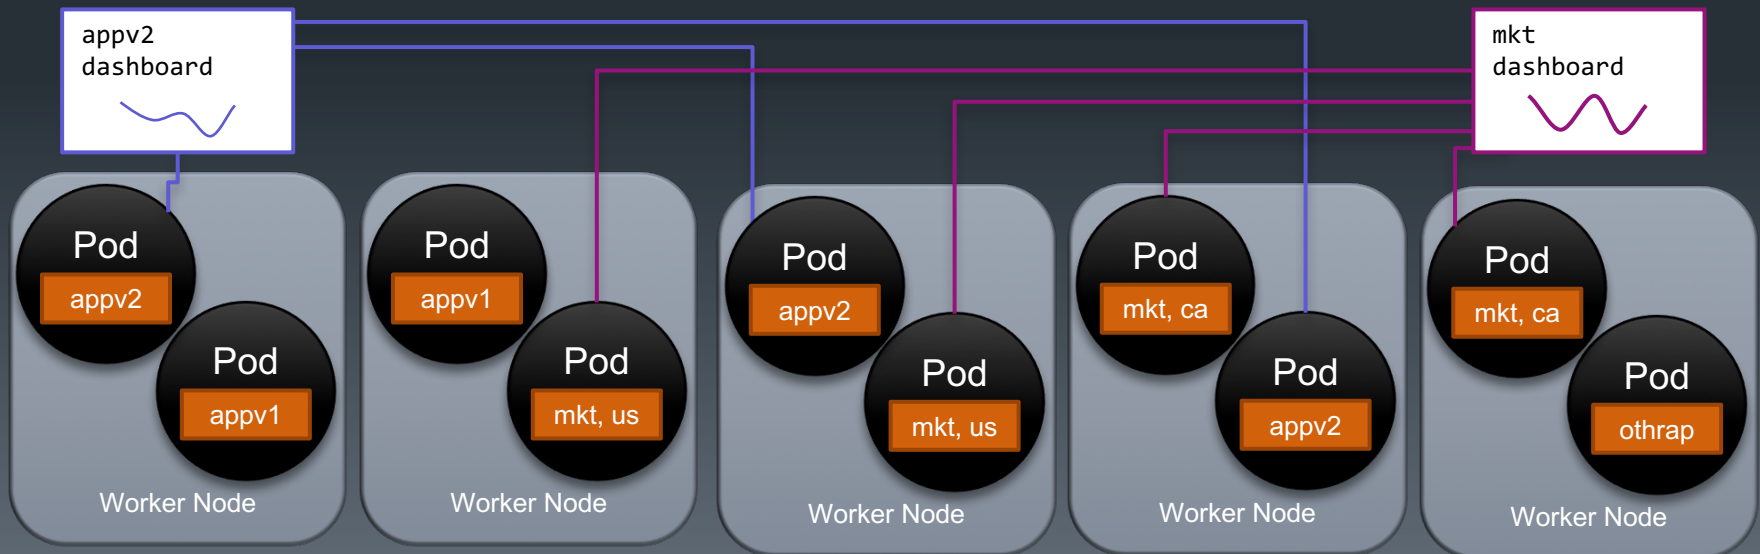
Controllers use Labels & Selectors to manage Pods

Labels

6

Copyright 2013-2017, RX-M LLC

- From one to many
 - When working on a single physical node, tools generally don't operate on containers in bulk
 - When moving to a cluster you want to easily scale out across nodes, requiring you to work in terms of **sets of things** instead of singletons
 - You want to keep those sets similarly configured
- In Kubernetes sets of pods are managed using two concepts:
 - **labels**
 - **Replica Sets**
- Every pod in Kubernetes has a set of key/value pairs associated with it called labels
- You can select a set of pods by constructing a query based on labels
- Kubernetes has no opinion on the "correct" way to organize pods
 - It is up to you to organize your pods in a way that makes sense to you
 - You can organize by application tiers, geographic location, development environment etc.
 - Labels are non-hierarchical, you can organize your pods in multiple ways simultaneously

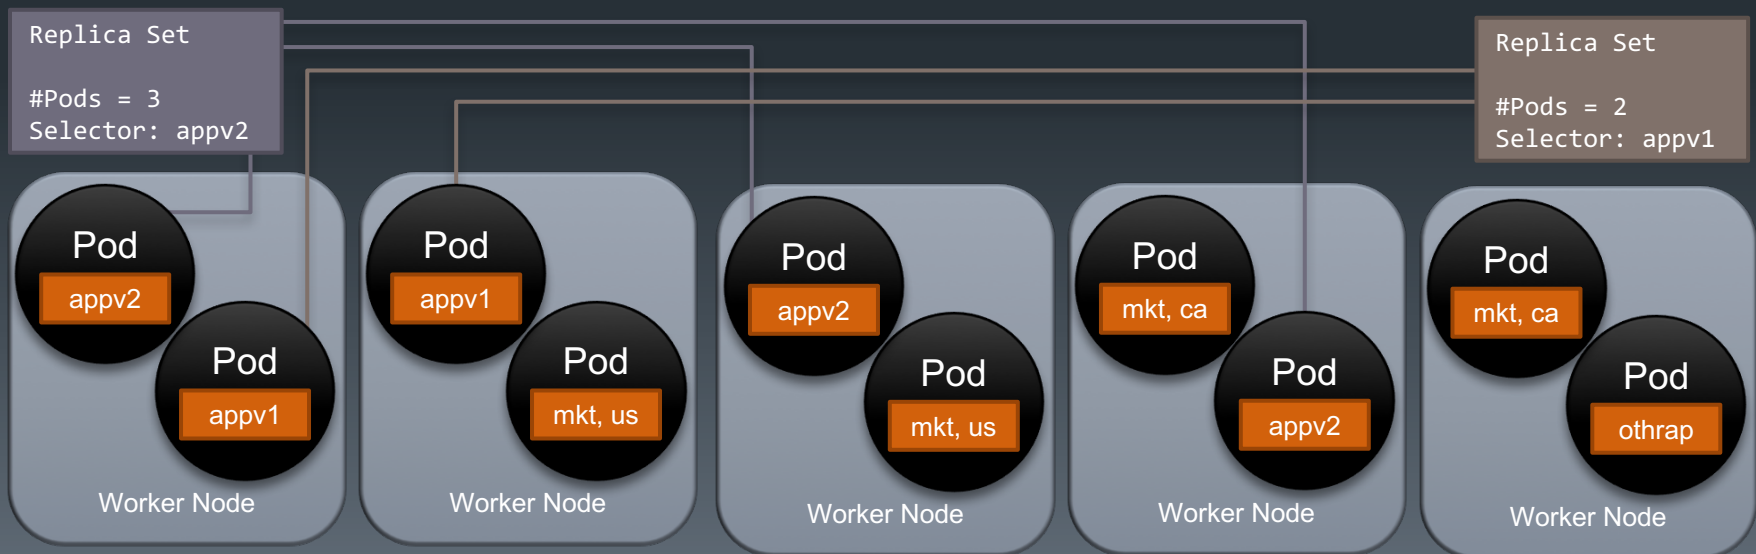


Replica Sets & Selectors

7

Copyright 2013-2017, RX-M LLC

- A Replica Set maintains a pool of pods based on a desired **replication count**, a pod **template** defined in a Deployment and a **label selector/query**
 - e.g. if you wanted to run a frontend tier with 3 pods, you would create a Deployment with an appropriate pod template (pointing at your container image) and a replicas count of 3
 - You would identify the set of pods that the Replica Set is managing with a label query, for example `env=prod,tier=fe`
- The Replica Set takes an easy to understand desired state and tirelessly works to make it true
- If you want to scale in or out all you have to do is change the desired replication count



Selectors

8

Copyright 2013-2017, RX-M LLC

- There are several options you can use to select on labels:
 - = or ==**
 - You can use either style to select keys with values equal to the string on the right
 - name = apache
 - !=**
 - Select keys with values that do not equal the string on the right
 - environment != test
 - in**
 - Select resources whose labels have keys with values in this set
 - tier in (web,app)
 - notin**
 - Select resources whose labels have keys with values not in this set
 - tier not in (lb,app)
 - <Key name>**
 - Use a key name only to select resources whose labels contain this key
 - tier
- Multiple selectors can be separated by commas and are logically “AND”ed
- Many kubectl subcommands offer the **-l** switch to specify label selectors (get, delete, etc.)

```
$ kubectl describe deployment logger
```

```
...
```

```
Selector:    deployment=demo,app=logger
```

```
...
```

```
$ kubectl get pod -l "deployment"
```

NAME	READY	STATUS	RESTARTS	AGE
logger-84ocd	1/1	Running	0	4m
logger-tc068	1/1	Running	0	4m
logger-zgtj9	1/1	Running	0	4m

```
$ kubectl get pod -l "deployment=demo"
```

NAME	READY	STATUS	RESTARTS	AGE
logger-84ocd	1/1	Running	0	4m
logger-tc068	1/1	Running	0	4m
logger-zgtj9	1/1	Running	0	4m

```
$ kubectl get pod -l "deployment=prod"
```

```
No resources found.
```

```
$ kubectl get pod -l "deployment!=prod"
```

NAME	READY	STATUS	RESTARTS	AGE
logger-84ocd	1/1	Running	0	5m
logger-tc068	1/1	Running	0	5m
logger-zgtj9	1/1	Running	0	5m

```
$ kubectl get pod -l "deployment in (prod,demo)"
```

NAME	READY	STATUS	RESTARTS	AGE
logger-84ocd	1/1	Running	0	4m
logger-tc068	1/1	Running	0	4m
logger-zgtj9	1/1	Running	0	4m

```
$ kubectl get pod -l "deployment in (prod,demo), app=logger"
```

NAME	READY	STATUS	RESTARTS	AGE
logger-84ocd	1/1	Running	0	9m
logger-tc068	1/1	Running	0	9m
logger-zgtj9	1/1	Running	0	9m

Deployment Config

9

Copyright 2013-2017, RX-M LLC

- The four first-level elements for a Deployment config are common among all top-level Kubernetes resource definitions
 - `kind` - in this case, set to Deployment
 - `apiVersion` - set to `apps/v1beta1` for the Deployment kind
 - `metadata` - labels
 - `spec` - the Deployment spec (see spec reference URL below)
- Keys related to the Deployment kind
 - `replicas`:
 - Defines the desired number of pods
 - `selector`:
 - Optional field that specifies a label selector for the Pods targeted by this deployment
 - Tells the Replica Set which pods to watch
 - If unspecified, `.spec.selector.matchLabels` defaults to `.spec.template.metadata.labels`
 - Assigning selectors that cause more than one RS to attempt to control a Pod is a **user error**
 - K8s tries to avoid oscillation but can not always
 - `template`:
 - Defines a template to launch a new pod
 - Contains the same elements defined for pod configs
- Note that, if used, selector values need to match the labels values specified in the pod template

```
apiVersion: apps/v1beta1
# before 1.6 use extensions/v1beta1
kind: Deployment
metadata:
  name: logger
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: logger
        deployment: demo
    spec:
      containers:
        - name: logger
          image: logagg:latest
          ports:
            - containerPort: 80
```

Deployment/Replication Resources

10

Copyright 2013-2017, RX-M LLC

API v1 [2015/7/10]

- Available since Kubernetes v0.x
- **kubectl rolling-update**
 - Imperative command used to migrate from one set of pod images to another
 - No support for rollback
- **ReplicationController**
 - Abbreviation: **rc**
 - Created by **kubectl run** in Kubernetes v1.0
 - Maintains a defined number of pod replicas
 - Spec includes:
 - replicas
 - selector
 - template (PodSpec) [required]
 - Will be deprecated in the future and ultimately removed

A Deployment that configures a ReplicaSet is now the recommended way to set up replication

Extension API v1Beta1 [2015/11/09]

- Available since Kubernetes v1.1
- **Deployment**
 - Abbreviation: **deploy**
 - Created by **kubectl run** in Kubernetes v1.1+
 - Declarative way to define the deployment of a set of pod images
 - Spec includes:
 - replicas
 - selector
 - template (PodSpec) [required]
 - Also:
 - strategy – how old pods will be replaced
 - minReadySeconds – how long after containers are up to wait before making Pod ready
 - revisionHistoryLimit – number of old ReplicaSets to retain (for rollback)
- **ReplicaSet**
 - Abbreviation: **rs**
 - Maintains a defined number of pod replicas
 - Used to perform a specific deployment

Rollouts

11

Copyright 2013-2017, RX-M LLC

- ReplicationController

- Rollouts

- `kubectl create -f abc.yaml`
- `kubectl rolling-update abc-v1 -f abc-v2.yaml`

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: website
  labels:
    bu: sales
spec:
  replicas: 3
template:
  metadata:
    labels:
      appname: webserver
      deployment: test
  spec:
    containers:
      - name: podweb
        image: nginx
        ports:
          - containerPort: 80
```

- Deployment

- Rollouts

- `kubectl create -f abc.yaml`
- `kubectl rollout pause deploy/abc`
- `kubectl rollout resume deploy/abc`
- `kubectl rollout status deploy/abc`
- `kubectl rollout history deploy/abc`
- `kubectl rollout undo deploy/abc`

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: website
  labels:
    bu: sales
spec:
  replicas: 3
template:
  metadata:
    labels:
      appname: webserver
      deployment: test
  spec:
    containers:
      - name: podweb
        image: nginx
        ports:
          - containerPort: 80
```

Simple Deployment / RS Walk Through

```
$ kubectl create -f mydep.yaml
deployment "logger" created
$ kubectl get deployment
NAME          DESIRED   CURRENT   UP-TO-DATE   AVAILABLE   AGE
logger        3         3         3            3           1m
$ kubectl get rs
NAME          DESIRED   CURRENT   READY   AGE
logger-1072152842  3         3         3       1m
$ kubectl get pod
NAME          READY   STATUS    RESTARTS   AGE
logger-1072152842-bj5bk  1/1     Running   0          2m
logger-1072152842-fh29f  1/1     Running   0          2m
logger-1072152842-q8cjc  1/1     Running   0          2m
$ kubectl describe deployment logger
Name:          logger
Namespace:     default
CreationTimestamp: Wed, 01 Aug 2017 12:48:13 -0800
Labels:        app=logger
                deployment=demo
Annotations:   deployment.kubernetes.io/revision=1
Selector:      app=logger,deployment=demo
Replicas:      3 updated | 3 total | 3 available | 0 unavailable
StrategyType:  RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=logger
            deployment=demo
  Containers:
    logger:
      Image:      logagg
      Port:       80/TCP
      Environment: <none>
      Mounts:      <none>
      Volumes:     <none>
Conditions:
  Type           Status      Reason
  ----           -
  Available      True        MinimumReplicasAvailable
  Progressing    True        NewReplicaSetAvailable
OldReplicaSets: <none>
NewReplicaSet: logger-1072152842 (3/3 replicas created)
Events:
  FirstSeen    LastSeen    Count   From          SubObjectPath  Type    Reason          Message
  -----
  3m           3m          1       {deployment-controller}                Normal    ScalingReplicaSet  Scaled up replica set logger-1072152842 to 3
```

```
$ cat mydep.yaml

apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: logger
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: logger
        deployment: demo
    spec:
      containers:
      - name: logger
        image: logagg:latest
        ports:
        - containerPort: 80
```

DaemonSets

- Ensures that **all nodes run a copy of a pod**
 - As nodes are added to the cluster, pods are added to them
- Good for running:
 - Storage daemons (e.g. glusterd, ceph, ...)
 - Log collectors (e.g. fluentd, logstash, ...)
 - Node monitoring (Prometheus Node Exporter, cAdvisor, collectd, New Relic agent, Ganglia gmond, ...)
- Marking a node unschedulable **does not affect DaemonSet** controller
- Must have a RestartPolicy equal to Always, or be unspecified, (which defaults to Always)
- Can make pods even when the scheduler has not been started
 - Also great for cluster bootstrap

```
apiVersion: extensions/v1beta1
kind: DaemonSet
metadata:
  name: monitor
  labels:
    app: monitor
spec:
  template:
    metadata:
      labels:
        app: monitor
    spec:
      containers:
        - name: monitor-container
          image: google/cadvisor
          ports:
            - containerPort: 8080
```

Jobs

14

Copyright 2013-2017, RX-M LLC

- Job
 - Creates one or more pods and ensures that a specified number of them successfully terminate
 - As pods successfully complete, the job tracks the successful completions; when a specified number of successful completions is reached, the job itself is complete
 - Will start a new Pod if the first pod fails or is deleted
 - 3 main types:
 - **Non-parallel** – one pod is started, job is complete as soon as Pod terminates successfully
 - **Parallel with fixed completion count**
 - Specify a non-zero positive value for `.spec.completions`
 - Job is complete when there is one successful pod for each value in the range: 1 to `.spec.completions`
 - **Parallel Jobs with a work queue** – support parallel processing of a set of independent but related work items
 - Examples: emails to be sent, frames to be rendered, files to be transcoded, ranges of keys in a NoSQL database to scan
 - Not designed to support closely-communicating parallel processes, as commonly found in scientific computing
 - In 1.5 moved to API version `batch/v1.Job` from `extensions/v1beta1.Job` (deprecated)

- CronJob

- Time-based Job resource: it **runs a job periodically on a given schedule**, written in Cron format
 - Point in time
 - Repeatedly at point in time
- Required **schedule** key - takes a Cron format string

CronJob requires a cluster at v1.5+ with batch/v2alpha1 API turned on by passing `--runtime-config=batch/v2alpha1` to the API server

```
apiVersion: batch/v2alpha1
kind: CronJob
metadata:
  name: hello
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: hello
              image: busybox
              args:
                - /bin/sh
                - -c
                - date; echo Hello from the Kubernetes cluster
          restartPolicy: OnFailure
```

Horizontal Pod Autoscaler

- Allows the number of pods in a replication controller or deployment to **scale automatically** based on metrics
- Metrics include
 - **Observed CPU utilization** – scales the number of pods in a Deployment, Replica Set or Replication Controller based on observed CPU utilization
 - Autoscaler periodically queries CPU utilization for the pods it targets
 - Default 30 seconds
 - Period controlled by `--horizontal-pod-autoscaler-sync-period`
 - Requires **Heapster** monitoring in your cluster for autoscaling to work
 - **Custom metrics** (alpha)
 - Makes use of the `autoscaling/v2alpha1` API
 - `targetCPUUtilizationPercentage` field replaced with an array called `metrics`
 - For example, queries per second or average request latency
 - Pods must have **cAdvisor**-specific custom metrics endpoint configured
- Launched via `kubectl autoscale` or YAML spec

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: apache
  namespace: default
spec:
  scaleTargetRef:
    apiVersion: apps/v1beta1
    kind: Deployment
    name: apache
  minReplicas: 1
  maxReplicas: 10
  targetCPUUtilizationPercentage: 50
```

Spec Policies

- It is a Kubernetes **best practice** to always use a Deployment with long running Pods
 - Even single pods with one container benefit from automatic restart by a Replica Set
- The config to the right demonstrates a JSON-based Deployment config
 - imagePullPolicy** is a container setting that tells Kubernetes when to pull the image for the container
 - Always (default if :latest tag is specified)
 - Never
 - IfNotPresent (default if :latest tag is not specified)
 - restartPolicy** is a Pod setting that applies to all container in the pod
 - Possible values:
 - Always (default)
 - OnFailure
 - Never
 - Failed containers are restarted by the Kubelet
 - Restarts use an exponential back-off delay in multiples of sync-frequency 0, 1x, 2x, 4x, 8x ... capped at 5 minutes and reset after 10 minutes of successful execution
 - A ReplicationController is only appropriate for pods with RestartPolicy = Always
 - dnsPolicy** sets the DNS search order with the Kubernetes cluster DNS server first followed by normal container DNS
 - More on cluster DNS in the services section

JSON Deployment example

```
{
  "apiVersion": "apps/v1beta1",
  "kind": "Deployment",
  "metadata": {
    "name": "frontend-controller",
    "labels": {
      "state": "serving"
    }
  },
  "spec": {
    "replicas": 2,
    "template": {
      "metadata": {
        "labels": {
          "app": "frontend"
        }
      },
      "spec": {
        "volumes": null,
        "containers": [
          {
            "name": "my-redis",
            "image": "redis",
            "ports": [
              {
                "containerPort": 80,
                "protocol": "TCP"
              }
            ],
            "imagePullPolicy": "IfNotPresent"
          }
        ],
        "restartPolicy": "Always",
        "dnsPolicy": "ClusterFirst"
      }
    }
  }
}
```


Health Checks

17

Copyright 2013-2017, RX-M LLC

- Kubernetes provides two layers of health checking
 - HTTP or TCP checks
 - K8s can attempt to connect to a particular endpoint and give a status of healthy on a successful connection
 - Application-specific health checks
 - K8s can be configured to use a command line scripts
- An RS/RC will restart a container failing a health check
- **livenessProbe**
 - Specifies a health check
 - Available options (typically only one of):
 - httpGet
 - Takes a path and port
 - Status codes between 200 and 399 are all considered healthy by the probe
 - tcpSocket
 - Takes a port
 - Any successful connect passes
 - exec
 - Takes a command
 - Exit code of 0 is success
- **readinessProbe**
 - Uses the same settings but uses test to assess container readiness

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
          livenessProbe: # don't use all at once!
            httpGet:
              path: /status/
              port: 80
            tcpSocket:
              port: 80
            exec:
              command: /usr/bin/health/ck.sh
          initialDelaySeconds: 30
          timeoutSeconds: 1
          periodSeconds: 10
```

Pre/Post Actions

18

Copyright 2013-2017, RX-M LLC

- Containers can be assigned lifecycle actions to run after the container starts and before the container stops
 - Convenient for initializing the started container or cleaning up before the container shutdown
- Management of the container blocks until lifecycle actions are complete
 - Unless the container process fails, in which case the action is aborted
- These hooks are called at least once (be prepared for more than one call!)
- postStart**
 - Called immediately after a container is created
 - If the handler fails, the container is terminated and restarted according to its restart policy
 - Other management of the container blocks until the action completes
- preStop**
 - Called immediately before a container is terminated
 - The container is terminated after the handler completes
 - The reason for termination is passed to the handler
 - "reason" values:
 - Delete - Delete command was issued via kubectl or the API
 - Health - Health check fails
 - Dependency - Dependency failure (e.g. disk mount failure)
 - Regardless of the outcome of the handler, the container is terminated
 - Other management of the container blocks until the hook completes

```
apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: apache-hook
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: apache-hook
    spec:
      containers:
        - name: apache-hook
          image: bitnami/apache:latest
          ports:
            - containerPort: 80
          lifecycle:
            postStart:
              httpGet:
                path: http://my.regsvr.com/register/
                port: 80
            preStop:
              exec:
                command: ["/usr/bin/apachectl","-k", "graceful- stop"]
```

Summary

- Controllers are designed to:
 - Run with a replication factor: Deployments & RSs or RCs
 - Run everywhere: DaemonSets
 - Run once: Jobs
 - Run on a schedule: CronJobs
 - Run stateful apps: StatefulSets (covered later)
- Deployments/ReplicaSets are key components in the orchestration of horizontally-scaled microservices
- RSs/RCs use kubelets to perform on-node monitoring and restart tasks
- Pods can be monitored for health and readiness
- Pods can be configured to perform post-start and pre-shutdown tasks

Lab 5

- Deployments and ReplicaSets

7: Managing State

Objectives

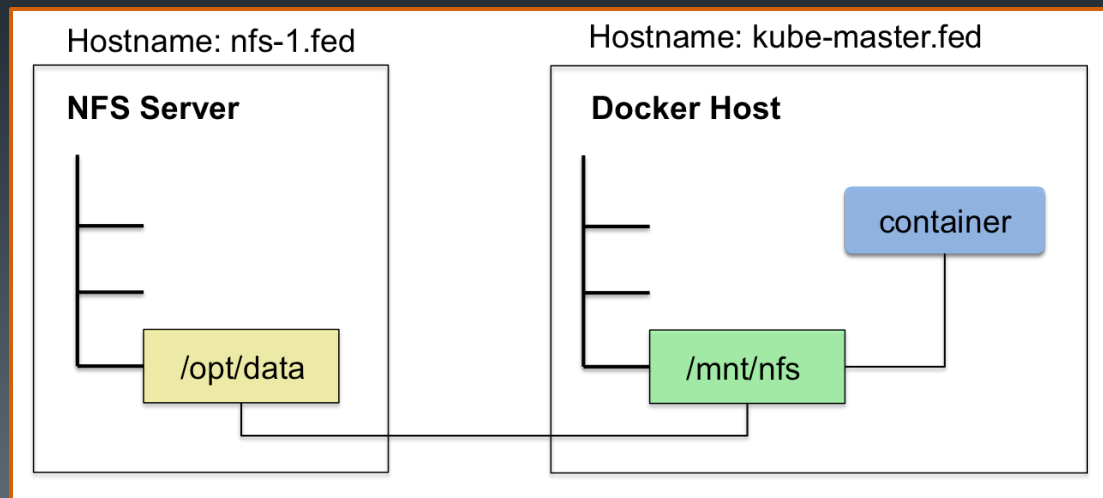
- Explore the lifecycle of volumes in Kubernetes
- Discover the different types of volumes available to Pods and containers in a Kubernetes cluster
- Examine ways to abstract storage providers
- Discuss creating and leveraging secrets in configs
 - Learn about encrypting secrets at rest

Kubernetes Volumes

23

Copyright 2013-2017, RX-M LLC

- A Volume is a container mountable filesystem feature with storage and persistence rules outside of the scope of the container
- Cases for volumes:
 - On-disk files in a container are **ephemeral**
 - When a container crashes, the kubelet will replace it by rerunning the original image, the files from the dead container **will be lost**
 - It may be necessary to persist data across container failures
 - When running containers together in a Pod it is often necessary to **share files** between those containers
 - Storage outside of the container mount namespace can be shared across containers
- A standard Kubernetes volume has a **lifetime the same as the pod** that encloses it
 - This is different from the Docker volume model, wherein volumes remain until explicitly deleted, regardless of whether there are any running containers using it
- It is important to remember that pods are anchored by the infrastructure container which cannot crash
 - Thus a pod volume outlives any containers that run within the Pod except the pause container preserving volume data across container restarts
 - Only when a Pod is deleted or the node the pod runs on crashes does the volume cease to exist



- Kubernetes supports many type of volumes, and a Pod can use any number of them simultaneously. Volume types supported include:
 - **emptyDir** – provides the pod with an empty directory (data is deleted when the Pod terminates)
 - **hostPath** – mounts a given host path into containers using the volume
 - Example: containerized kubelet that needs `/var/lib/docker` mounted to talk to `docker.sock`
 - **gcePersistentDisk** – mounts a Google Compute Engine block storage volume
 - **awsElasticBlockStore** – mounts an Amazon Web Service EBS volume
 - **nfs** – mounts an NFS share
 - **iscsi** – allows an existing iSCSI (SCSI over IP) volume to be mounted
 - **fc (fibre channel)** – allows an existing fibre channel volume to be mounted in a pod
 - **flocker** – mounts a flocker volume
 - **glusterfs** – mounts a glusterfs volume (Glusterfs is an open source networked filesystem)
 - **rbd** – mounts a Rados Block Device volume
 - **cephfs** – mounts an existing CephFS volume
 - **gitRepo** – mounts an empty directory and clones a git repository into it for your pod to use
 - **secret** – mounts a volume as an in-memory tmpfs (never touches the host disk)
 - Used to pass sensitive information, such as passwords, to pods
 - **downwardAPI** – mounts pod metadata
 - **projected** – maps several existing volume sources into the same directory
 - Following types of volume sources can be projected: `secret`, `downwardAPI`, `configMap`
 - **azureFileVolume** – mounts an Azure File Volume
 - **azureDisk** – mounts a Microsoft Azure Data Disk
 - **vsphereVolume** – mounts a vSphere VMDK Volume
 - **Quobyte** – allows an existing Quobyte volume to be mounted
 - **PortworxVolume** – elastic block storage layer that runs hyperconverged with Kubernetes
 - **ScaleIO** – software-based storage platform of scalable shared block networked storage
 - **StorageOS** – allows an existing StorageOS volume to be mounted into your pod
 - **PersistentVolume (PV)** – storage provisioned by an administrator aimed at abstracting storage drivers
 - **local** (alpha in 1.7) – local storage device such as a disk, partition or directory
 - Can only be used as a statically created PersistentVolume

emptyDir

- One of the easiest ways to achieve improved persistence amid container crashes and data sharing within a pod is to use the emptydir volume
- Can be mounted at the **same or different paths** in each container
- Created when a Pod is assigned to a Node, **erased when a Pod is removed**
 - Persists beyond container crashes
- Stored on the medium backing the Node (SSD, network storage, etc.)
 - Specifying Memory creates a RAM-backed filesystem
 - Cleared on machine reboot
 - Counts against container memory limits

```
apiVersion: v1
kind: Pod
metadata:
  name: memory-pd-pod
spec:
  containers:
  - image: nginx:1.13
    ports:
    - containerPort: 80
    name: memory-pd-nginx
    volumeMounts:
    - mountPath: /memory-pd
      name: memory-volume
  volumes:
  - name: memory-volume
    emptyDir:
      medium: Memory
```

hostPath

- Mounts a file or directory from the Node's filesystem into a Pod
- Pods with identical configs may behave differently on different Nodes due to differences in files on the Nodes
- Resource-aware scheduling will **not account for resources** used by hostPath
- Created directories only writeable by root
 - Config options:
 - Run process as root in privileged container
 - Modify file perms on the host to write to the volume
- Unlike emptyDir, which is erased when a Pod is removed, the contents of hostPath volumes are preserved and the volume is merely unmounted

```
apiVersion: extensions/v1beta1
kind: DaemonSet
metadata:
  labels:
    app: cadvisor
spec:
  template:
    metadata:
      labels:
        app: cadvisor
    spec:
      containers:
        - name: cadvisor
          image: google/cadvisor
          securityContext:
            privileged: true
          volumeMounts:
            - name: var-run
              mountPath: /var/run
              readOnly: false
            - name: sys
              mountPath: /sys
              readOnly: true
            - name: docker
              mountPath: /var/lib/docker
              readOnly: true
      volumes:
        - name: var-run
          hostPath:
            path: /var/run
        - name: sys
          hostPath:
            path: /sys
        - name: docker
          hostPath:
            path: /var/lib/docker
```

Persistent Disks

27

- Types:
 - nfs**
 - NFS can be mounted by multiple writers simultaneously
 - iscsi**
 - Can be mounted as read-only by multiple consumers simultaneously
 - No simultaneous writers allowed
 - fc (fibre channel)**
 - flocker**
 - Flocker is an open-source clustered container data volume manager that provides orchestration of data volumes backed by a variety of storage backends
 - glusterfs**
 - GlusterFS can be mounted by multiple writers simultaneously
 - rbd**
 - RBD is can be mounted as read-only by multiple consumers simultaneously but can only be mounted by a single consumer in read-write mode - no simultaneous writers allowed
 - cephfs**
 - CephFS can be mounted by multiple writers simultaneously
- A persistent volume that exists outside the container allows one to save important data across containers outages
- Volumes at the pod level can be shared between containers in the same application stack and within the same pod
- Unlike emptyDir, which is erased when a Pod is removed, the contents of nfs, iscsi, fc, flocker, glusterfs, rbd and cephfs volumes are preserved and the volume is merely unmounted

```
apiVersion: v1
kind: Pod
metadata:
  name: flocker-pod
spec:
  containers:
    - name: web
      image: nginx
      ports:
        - name: web
          containerPort: 80
      volumeMounts:
        - name: www-root
          mountPath: /usr/share/nginx/html
  volumes:
    - name: www-root
      flocker:
        datasetName: my-flocker-vol

---
apiVersion: v1
kind: Pod
metadata:
  name: fc-pod
spec:
  containers:
    - image: mysql
      name: fc-pod
      volumeMounts:
        - name: fc-vol
          mountPath: /var/lib/mysql
  volumes:
    - name: fc-vol
      fc:
        targetWWNs: ['500a0982991b8dc5', '500a0982891b8dc5']
        lun: 2
        fsType: ext4
        readOnly: true
```

Cloud Volumes

- Cloud vendors (particularly AWS and Google) offering block devices can support named volume mounting
- gcePersistentDisk** – volume mounts a Google Compute Engine (GCE) Persistent Disk into your pod
 - Restrictions:
 - Nodes using GCE-PDs must be GCE VMs
 - GCE-PDs need to be in the same **project** and **zone** as the cluster
 - You can attach **up to 16 persistent disks** to most instances and up to 64 TB of total persistent disk space
 - Instances with shared-core machine types or custom machine types with less than 3.75 GB of memory are limited to a maximum of 4 persistent disks and 3 TB of total persistent disk space
 - You can attach up to 128 persistent disks to instances with predefined machine types depending on the number of vCPUs in that instance (Beta)
- awsElasticBlockStore** – volume mounts an AWS EBS Volume into your pod
 - Restrictions
 - Nodes using EBS volumes must be AWS EC2 instances
 - Instances must be in the same **region** and **availability-zone** as the EBS volume
 - EBS only supports a **single EC2 instance mounting a volume**
 - Amazon recommends a **maximum of 40 EBS volumes** per Node with one of those 40 reserved for the root volume
 - Attaching more than 40 volumes can cause boot failures
- You must create a GCE-PD or EBS volume **before** you can use it
- Unlike emptyDir, which is erased when a Pod is removed, the contents of GCE-PD and EBS volumes are preserved and the volume is merely **unmounted**

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    cloud: test-gce
spec:
  containers:
    - image: nginx:1.9.1
      name: gce-nginx
      volumeMounts:
        - mountPath: /usr/share/nginx/html
          name: gce-pd
  volumes:
    - name: gce-pd
      gcePersistentDisk:
        pdName: mysite-volume-1
        fsType: ext4
---
apiVersion: v1
kind: Pod
metadata:
  labels:
    cloud: test-aws
spec:
  containers:
    - image: nginx:1.13
      name: aws-nginx
      volumeMounts:
        - mountPath: /usr/share/nginx/html
          name: awsvol
    - name: awsvol
      awsElasticBlockStore:
        volumeID: vol-066c9038349ed7fdd
        fsType: ext4
```

Data Provider Abstraction

- Abstracts details of how storage is provided from how it is consumed
 - **PersistentVolume** (PV) – piece of storage in the cluster that has been provisioned dynamically or by a cluster admin and is available for consumption
 - Are volume plugins like Volumes but have a lifecycle independent of any individual pod that uses the PV
 - Capture the details of the implementation of the storage: NFS, iSCSI, or cloud-provider-specific storage, etc
 - **persistentVolumeClaim** (PVC) – PVCs consume PV resources the way Pods consume node resources
 - Claims can request storage sizes and access modes
 - Like Pods, which can request specific levels of resources (CPU and Memory)
 - Pods access storage by using a PVC as a volume
 - **StorageClass** provides a way for administrators to describe the “classes” of storage they offer
 - Enables a variety of PVs that differ in more ways than just size and access modes, without exposing users to the details of how those volumes are implemented
 - Map to quality-of-service levels, backup policies, or to arbitrary policies determined by cluster admins
 - Kubernetes is unopinionated about what classes represent

```
apiVersion: v1
kind: Pod
metadata:
  name: frontendpod
spec:
  containers:
    - name: myfrontend
      image: nginx
      volumeMounts:
        - mountPath: "/var/www/html"
          name: fivegigpvc
  volumes:
    - name: fivegigpvc
      persistentVolumeClaim:
        claimName: fivegigslowpvc
```

Persistent Volumes

30

Copyright 2013-2017, RX-M LLC

- Provisioning
 - **Static** – cluster admin creates a number of PVs
 - Include the details of the real storage
 - Exist as resources in the API
 - **Dynamic** – PersistentVolume Provisioner may try to provision a PV when none of the static PVs matches a user's PVC
 - Based on Storage Classes and volume plugin feature availability
- Capacity – PVs will have a storage capacity set using the PV's capacity attribute
- Access Modes – can be set to modes supported by a given provider
 - **ReadWriteOnce** (RWO) – can be mounted as read-write by a single node
 - **ReadOnlyMany** (ROX) – can be mounted read-only by many nodes
 - **ReadWriteMany** (RWX) – can be mounted as read-write by many nodes
 - PV can only be mounted using one access mode at a time, even if it supports many
 - ex: GCE-PD can be mounted as ReadWriteOnce by a single node or ReadOnlyMany by many nodes but not at the same time
- Reclaim policy for a PV tells the cluster what to do with the volume after it has been released of its claim, can be:
 - **Retained** – requires manual reclamation
 - Previous claimant's data remains on the volume until cleaned or the PV is deleted
 - **Recycled** - performs a scrub (`rm -rf /thevolume/*`) on the volume and makes it available again for a new claim
 - Must use a volume plugin that supports the recycled feature
 - **Deleted** – removes both the PV object as well as deleting the associated external storage, such as an AWS EBS, GCE PD, Azure Disk, or Cinder volume
 - Dynamically provisioned PVs are always deleted
- Class (optional) – ensures that the PV will only be bound to PVCs requesting the given Storage Class

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: fivegigslowpv
spec:
  capacity:
    storage: 5Gi
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Recycle
  storageClassName: slow
  nfs:
    path: /tmp
    server: 172.17.0.2
```

Persistent Volume Claims

31

Copyright 2013-2017, RX-M LLC

- Created with a specific amount of storage and an access mode
 - Resources – claims request quantities of storage
 - Access modes – PVCs use the same conventions as PVs
 - Mounting claims with "Many" modes (ROX, RWX) only possible within one namespace
- Class (optional) – only PVs of the requested class can be bound to the PVC
- Selector (optional) – enables granular filtering of available PVs
 - Only volumes with labels that match the selector can be bound to the claim
 - `matchLabels` – volume must have a label with this value
 - `matchExpressions` – list of requirements made by specifying key, a list of values, and an operator that relates the key and values
 - Valid operators: In, NotIn, Exists, DoesNotExist
 - Requirements are ANDed together—all must be satisfied to match
- Control loop matches new PVCs with PVs
 - A dynamically-provisioned PV will always bind to the PVC that requested it
- Will remain unbound indefinitely if a matching PV does not exist
 - Claims are bound as matching volumes become available
 - Cluster with many 50Gi PVs would not match any to a 100Gi PVC

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: fivegigslovpvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 5Gi
  storageClassName: slow
  selector:
    matchLabels:
      release: "stable"
    matchExpressions:
      - {key: environment, operator: In, values: [dev]}
```

Storage Classes

32

Copyright 2013-2017, RX-M LLC

- Admins set the name and other parameters of a class when first creating StorageClass objects
 - Cannot be updated once they are created
- **Provisioner** – determines what volume plugin is used for provisioning PVs
 - ex: GCEPersistentDisk, AWSElasticBlockStore, VsphereVolume, Flocker, etc.
 - **Internal** provisioners prefixed with `kubernetes.io` and shipped with Kubernetes
 - **External** provisioners are independent programs that follow a specification defined by Kubernetes
 - ex: NFS doesn't provide an internal provisioner, but an external provisioner can be used
- **Parameters** – describe attributes of volumes belonging to the storage class
 - Are provisioner specific though some may be similar
 - Best to examine docs of a given provisioner for details on parameters
- Default StorageClass can be specified for PVCs that don't request any particular class to bind to

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: slow
provisioner: kubernetes.io/gce-pd
parameters:
  type: pd-standard
  zones: us-central1-a, us-central1-b

---
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: standard
provisioner: kubernetes.io/aws-efs
parameters:
  type: gp2
```


StatefulSets

- Beta as of 1.5, replacement for PetSets
- Pods in a StatefulSet have a **unique ordinal index**, a **stable network identity**, **stable storage**
 - Ordered deployment, scaling, deletion and termination
 - For a StatefulSet with N replicas, Pods are given a unique index, in order from {0...N-1}
 - K8s v1.7+ uses podManagementPolicy to determine ordering
 - OrderedReady – the first pod must be “running and ready” prior to the following pods start (Also terminated in reverse order)
 - Parallel – launch or terminate all Pods in parallel
 - Each Pod in the StatefulSet will be assigned an integer ordinal that is unique over the Set
 - Unique network identifiers – each Pod in a StatefulSet derives its hostname from the name of the StatefulSet and the ordinal of the Pod
 - Requires the use of a headless service
 - The CNAME of the headless service points to SRV records (one for each Pod)
 - The IP can change so clients should use the SRV
 - kube-proxy never acts on headless services
 - Stable persistent storage – one PV for each VolumeClaimTemplate
 - PVs associated with the Pods’ PVCs are not deleted when the Pods, or StatefulSet are deleted & must be done manually
- Update strategies (K8s v1.7+) – configure automated rolling updates
 - **OnDelete** (default) – users must manually delete Pods to cause the controller to create new Pods that reflect modifications made to a StatefulSet’s .spec.template
 - **RollingUpdate** – deletes and recreates each Pod in the same order as Pod termination updating each Pod one at a time
 - It will wait until an updated Pod is Running and Ready prior to updating its predecessor
 - Can be partitioned, by specifying a .spec.updateStrategy.rollingUpdate.partition
 - If a partition is specified, all Pods with an ordinal that is greater than or equal to the partition will be updated
 - All Pods with an ordinal that is less than the partition will not be updated
 - If deleted, will be recreated at the previous version
 - Useful if you want to stage an update, roll out a canary, or perform a phased roll out

```
# Headless service
apiVersion: v1
kind: Service
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  ports:
    - port: 80
  clusterIP: None
  selector:
    app: nginx

---
apiVersion: apps/v1beta1
kind: StatefulSet
metadata:
  name: web
spec:
  serviceName: nginx
  replicas: 3
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
          volumeMounts:
            - name: www
              mountPath: /usr/share/nginx/html
  volumeClaimTemplates:
    - metadata:
        name: www
      spec:
        accessModes: [ "ReadWriteOnce" ]
        storageClassName: slow
        resources:
          requests:
            storage: 5Gi
```

Secrets

34

Copyright 2013-2017, RX-M LLC

- Kubernetes objects used to hold sensitive information
 - Passwords
 - OAuth tokens
 - ssh keys
- Safer and more flexible than putting it verbatim in a pod definition or in a Docker image
- Users can create secrets
 - The system also creates some secrets
- Can only be referenced by pods in the **same namespace** where created
- To use a secret, a pod needs to reference the secret
- Each item must be **base64 encoded**
- Limited to 1MB in size
 - Discourages large secrets which would exhaust apiserver and kubelet memory
- Stored as **plaintext** in etcd
 - Secret data is at rest on the disk that etcd uses

```
$ echo -n "admin" | base64
YWRtaW4=
```

```
$ echo -n "1f2d1e2e67df" | base64
MWYyZDF1MmU2N2Rm
```

```
apiVersion: v1
kind: Secret
metadata:
  name: prod-db-secret
type: Opaque
data:
  password: MWYyZDF1MmU2N2Rm
  username: YWRtaW4=
```

Mounting Secrets

- A secret can be used with a Pod in three ways:
 - As files in a volume
 - Environment variables
 - Used by kubelet when pulling images for the pod
- When a pod is created secret existence is not checked
- If the kubelet can not fetch the secret:
 - It will report a pod event
 - The kubelet will periodically retry fetching the secret
 - Once fetched the kubelet will create and mount the secret volume
 - None of the pod's containers will start until all the pod's volumes are mounted
- Files
 - You can specify the permission mode bits a secret or files of part of a secret will have
 - If you don't specify any, 0644 is used by default
 - When a secret being consumed in a volume is updated, projected keys are **eventually updated** as well
 - Depends on kubelet sync period
- Environment variables
 - Appear as normal environment variables containing the base-64 decoded values of the secret data
- An **imagePullSecret** is a way to pass a secret that contains a Docker (or other) image registry password to Kubelet so it can pull a private image on behalf of your Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: secret-file-pod
  labels:
    secretpod: db-client
spec:
  volumes:
    - name: secret-volume
      secret:
        secretName: prod-db-secret
  containers:
    - name: secret-file-container
      image: redis
      volumeMounts:
        - name: secret-volume
          readOnly: true
          mountPath: "/etc/secret-volume"
---
apiVersion: v1
kind: Pod
metadata:
  name: secret-env-pod
spec:
  containers:
    - name: secret-env-container
      image: redis
      env:
        - name: SECRET_USERNAME
          valueFrom:
            secretKeyRef:
              name: prod-db-secret
              key: username
        - name: SECRET_PASSWORD
          valueFrom:
            secretKeyRef:
              name: prod-db-secret
              key: password
```

Encrypted Secret Data

36

Copyright 2013-2017, RX-M LLC

- Encryption at rest is **alpha** as of version 1.7.0
 - Set `--experimental-encryption-provider-config` flag on the apiserver to point to the location of the EncryptionConfig file
 - After restarting apiserver, newly created or updated secrets should be **encrypted when stored**
- EncryptionConfig file contains keys that can decrypt content in etcd
 - Restrict permissions on masters so only the user who runs apiserver can read it
- The **providers** array is an ordered list of the possible encryption providers:
 - **aescbc** – recommended, strongest but slowest
 - AES-CBC PKCS#7 padding
 - **secretbox** – faster than, but not as strong as aescbc
 - XSalsa20 and Poly 1305
 - **aesgcm** – fastest, not recommended if not using automated key rotation as must be rotated every 200k writes
 - AES-GCM with random nonce
 - **identity** – no encryption
 - When set as first provider, will disable encryption at rest
- Provider key secrets generated with a 32 byte random key and base64 encoding it
 - ex: `head -c 32 /dev/urandom | base64`

```
kind: EncryptionConfig
apiVersion: v1
resources:
- resources:
  - secrets
  providers:
  - aesgcm:
    keys:
    - name: key1
      secret: c2Vjc mV0IG1zIHN1Y3VyZQ==
    - name: key2
      secret: dGhpcyBpcyBwYXNzd29yZA==
  - aescbc:
    keys:
    - name: key1
      secret: c2Vjc mV0IG1zIHN1Y3VyZQ==
    - name: key2
      secret: dGhpcyBpcyBwYXNzd29yZA==
  - secretbox:
    keys:
    - name: key1
      secret: YWJjZGVmZ2hpamtsbW5vcHFyc3R1dnd4eXoxMjM0NTY=
  - identity: {}
```

- Many applications require configuration via some combination of config files, command line arguments, and environment variables
- These configuration **artifacts should be decoupled** from image content in order to keep containerized applications portable
- The ConfigMap resource provides mechanisms to **inject containers with configuration data** while keeping containers agnostic of Kubernetes
- ConfigMap can be used to store fine-grained information like individual properties or coarse-grained information like entire config files or JSON blobs

```
# ConfigMap of prometheus.yml config file
```

```
apiVersion: v1
```

```
kind: ConfigMap
```

```
metadata:
```

```
  name: prometheus-config
```

```
data: prometheus.yml: |
```

```
  global:
```

```
    scrape_interval: 30s
```

```
    scrape_timeout: 30s
```

```
  scrape_configs:
```

```
- job_name: 'kubernetes-apisservers'
```

```
  kubernetes_sd_configs:
```

```
- role: endpoints
```

```
  scheme: https
```

```
  tls_config:
```

```
    ca_file: /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
```

```
  bearer_token_file: /var/run/secrets/kubernetes.io/serviceaccount/token
```

```
  relabel_configs:
```

```
- source_labels: [__meta_kubernetes_namespace, __meta_kubernetes_service_name]
```

```
  action: keep
```

```
  regex: default;kubernetes;https
```

```
- job_name: 'kubernetes-nodes'
```

```
...
```

```
# prometheus.yml
```

```
# Example scrape configuration for running Prometheus on a Kubernetes cluster.
```

```
## Kubernetes Labels will be added as Prometheus Labels on metrics via the
```

```
# `labelmap` relabeling action.
```

```
scrape_configs:
```

```
- job_name: 'kubernetes-apisservers'
```

```
  kubernetes_sd_configs:
```

```
- role: endpoints
```

```
  scheme: https
```

```
  tls_config:
```

```
    ca_file: /var/run/secrets/kubernetes.io/serviceaccount/ca.crt
```

```
    # insecure_skip_verify: true
```

```
  bearer_token_file: /var/run/secrets/kubernetes.io/serviceaccount/token
```

```
  relabel_configs:
```

```
- source_labels: [__meta_kubernetes_namespace, __meta_kubernetes_service_name]
```

```
  action: keep
```

```
  regex: default;kubernetes;https
```

```
- job_name: 'kubernetes-nodes'
```

```
...
```

ConfigMaps vs Secrets

Paul Morie (author of both resources):

“

I'm the author of both of these features. The idea is that you should:

1. Use secrets for things which are actually secret like API keys, credentials, etc
2. Use config map for not-secret configuration data

In the future there will likely be some differentiators for secrets like rotation or support for backing the secret API w/ HSMs, etc. In general we like intent-based APIs, and the intent is definitely different for secret data vs. plain old configs.

Hope that helps.

“

<http://stackoverflow.com/questions/36912372/kubernetes-secrets-vs-configmaps>



Summary

- Kubernetes volumes have the same lifetime as the pods that declare them
 - Kubernetes volumes are not exactly the same as Docker volumes
- Kubernetes supports many volume types & providers
- Volumes are enabled via plugins
- Kubernetes provides powerful storage abstraction using PVs, PVCs & StorageClasses
- StatefulSets leverage abstracted storage for stateful applications
- Secrets can be used to pass sensitive data to Pods
 - kubelet can use secrets for logging into registries
 - Secrets can be encrypted by using EncryptionConfigs

Lab 7

- Volumes