

Kubernetes

Lab 5 – Deployments and Replica Sets

In this lab we will explore the nature of Kubernetes deployments and replica sets and how to work with them.

Deployments

A deployment provides declarative updates for pods and replica sets. You describe the desired state in a deployment object, and the deployment controller will change the actual state to the desired state at a controlled rate for you. You can define deployments to create new resources, or replace existing ones by new ones. Typical uses:

- bring up a replica set and (indirectly) its pods
- capturing the results and status of a deployment
- updating an existing deployment to recreate pods with a new image (rolling updates)
- rolling back to an earlier deployment revision if the current deployment isn't stable
- pausing and resuming a deployment

ReplicaSets

Replica sets (RS) supersede the older replication controller (RC) resource type. Replica sets support the set-based selectors as well as equality-based selector requirements (RCs only supported equality.) While replica sets can be used independently, they are mainly used by deployments as a mechanism to orchestrate pod creation, deletion, and updates. When you use deployments you don't have to worry about managing the replica sets that they create; deployments own and manage their replica sets.

ReplicaSets ensure that a specified number of pod "replicas" are running at all times. If there are too many, it will kill some. If there are too few, it will start more. Unlike in the case where a user directly created pods, a ReplicaSet replaces pods that are deleted or terminated for any reason, such as in the case of node failure or disruptive node maintenance (e.g. a kernel upgrade, etc.)

For this reason the Kubernetes team recommends that you use a Deployment/ReplicaSet even if your application requires only a single pod. ReplicaSets are like

process supervisors in many ways but monitor processes on multiple nodes at once. A ReplicaSet delegates local container restarts to some agent on the node (e.g., Kubelet or Docker.)

A ReplicaSet is only appropriate for pods with *RestartPolicy = Always* (if the RestartPolicy is not set, the default value is *Always*.) A ReplicaSet will refuse to instantiate any pod that has a different restart policy.

A ReplicaSet will never terminate on its own, but it isn't expected to be as long-lived as services. Services may be composed of pods controlled by multiple ReplicaSets, and it is expected that many ReplicaSets may be created and destroyed over the lifetime of a service (for instance, to perform an update of pods that run the service.) Both services themselves and their clients should remain oblivious to the ReplicaSets that maintain the pods of the services.

Now to create some Deployments/ReplicaSets.

1. A Simple Deployment

As a first exploratory step lets create a simple deployment which stands up three nginx pods. Create a config file similar to the following to accomplish this task:

```
user@ubuntu:~$ cd ~
user@ubuntu:~$ mkdir dep
user@ubuntu:~$ cd dep
user@ubuntu:~/dep$
user@ubuntu:~/dep$ vim mydep.yaml
user@ubuntu:~/dep$ cat mydep.yaml

apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: website
  labels:
    bu: sales
spec:
  replicas: 3
  template:
    metadata:
      labels:
        appname: webserver
        targetenv: demo
```

```
spec:
  containers:
  - name: podweb
    image: nginx:1.7.9
    ports:
    - containerPort: 80
```

The first thing you might notice is that Deployments are a beta resource type. Deployments were added to Kubernetes 1.2 and are the go forward solution for deploying replicated pods. The spec for Replication Controllers (part of the v1 API) is almost the same as the spec for Deployments though deployments add a few key features such as the ability to specify upgrades declaratively. The beta specification for Deployments can be found [here](https://kubernetes.io/docs/api-reference/v1.7/#deployment-v1beta1-apps):

<https://kubernetes.io/docs/api-reference/v1.7/#deployment-v1beta1-apps>

Now create the Deployment using the `kubectl create` subcommand and verify that the Deployment, its ReplicaSet and pods are up with the `get` subcommand:

```
user@ubuntu:~/dep$ kubectl create -f mydep.yaml

deployment "website" created

user@ubuntu:~/dep$
```

```
user@ubuntu:~/dep$ kubectl get deploy,rs,pods
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deploy/website	3	3	3	3	27s

NAME	DESIRED	CURRENT	READY	AGE
rs/website-205870431	3	3	3	27s

NAME	READY	STATUS	RESTARTS	AGE
po/website-205870431-7cd9k	1/1	Running	0	27s
po/website-205870431-h530h	1/1	Running	0	27s
po/website-205870431-l7czq	1/1	Running	0	27s

```
user@ubuntu:~/dep$
```

While everything appears to be running we can verify that there are no scheduling cycles or fail/restart activities by examining the system events. We have viewed

resource specific events in the past using the `kubectl describe` subcommand. This time we'll use the `kubectl get events` subcommand to view cluster wide events:

```
user@ubuntu:~/dep$ kubectl get events | grep website
```

1m	1m	1	website-205870431-7cd9k	Pod		Normal	
Scheduled			default-scheduler	Successfully assigned website-205870431-7cd9k to ubuntu			
1m	1m	1	website-205870431-7cd9k	Pod	spec.containers{podweb}	Normal	
Pulling			kubelet, ubuntu	pulling image "nginx:1.7.9"			
1m	1m	1	website-205870431-7cd9k	Pod	spec.containers{podweb}	Normal	Pulled
kubelet, ubuntu				Successfully pulled image "nginx:1.7.9"			
1m	1m	1	website-205870431-7cd9k	Pod	spec.containers{podweb}	Normal	
Created			kubelet, ubuntu	Created container with id 15a14aa6f5faa02c20c65406572a07073ff454ff30a6a65ad8e3baa46be2b514			
1m	1m	1	website-205870431-7cd9k	Pod	spec.containers{podweb}	Normal	
Started			kubelet, ubuntu	Started container with id 15a14aa6f5faa02c20c65406572a07073ff454ff30a6a65ad8e3baa46be2b514			
1m	1m	1	website-205870431-h530h	Pod		Normal	
Scheduled			default-scheduler	Successfully assigned website-205870431-h530h to ubuntu			
1m	1m	1	website-205870431-h530h	Pod	spec.containers{podweb}	Normal	
Pulling			kubelet, ubuntu	pulling image "nginx:1.7.9"			
1m	1m	1	website-205870431-h530h	Pod	spec.containers{podweb}	Normal	Pulled
kubelet, ubuntu				Successfully pulled image "nginx:1.7.9"			
1m	1m	1	website-205870431-h530h	Pod	spec.containers{podweb}	Normal	
Created			kubelet, ubuntu	Created container with id 771547aa87b75b6c9debb86a43c1676a09737afefe1c1d13fe8f52d2b6769c1e			
1m	1m	1	website-205870431-h530h	Pod	spec.containers{podweb}	Normal	
Started			kubelet, ubuntu	Started container with id 771547aa87b75b6c9debb86a43c1676a09737afefe1c1d13fe8f52d2b6769c1e			
1m	1m	1	website-205870431-l7czq	Pod		Normal	
Scheduled			default-scheduler	Successfully assigned website-205870431-l7czq to ubuntu			
1m	1m	1	website-205870431-l7czq	Pod	spec.containers{podweb}	Normal	
Pulling			kubelet, ubuntu	pulling image "nginx:1.7.9"			
1m	1m	1	website-205870431-l7czq	Pod	spec.containers{podweb}	Normal	Pulled
kubelet, ubuntu				Successfully pulled image "nginx:1.7.9"			
1m	1m	1	website-205870431-l7czq	Pod	spec.containers{podweb}	Normal	
Created			kubelet, ubuntu	Created container with id 294f3a303d660a4771e0f912a21d9758ee62f6b74bcf042445dc75838644d65f			
1m	1m	1	website-205870431-l7czq	Pod	spec.containers{podweb}	Normal	
Started			kubelet, ubuntu	Started container with id 294f3a303d660a4771e0f912a21d9758ee62f6b74bcf042445dc75838644d65f			
1m	1m	1	website-205870431	ReplicaSet		Normal	

```

SuccessfulCreate    replicaset-controller    Created pod: website-205870431-h530h
1m                  1m                      1      website-205870431      ReplicaSet              Normal
SuccessfulCreate    replicaset-controller    Created pod: website-205870431-7cd9k
1m                  1m                      1      website-205870431      ReplicaSet              Normal
SuccessfulCreate    replicaset-controller    Created pod: website-205870431-l7czq
1m                  1m                      1      website                Deployment              Normal
ScalingReplicaSet    deployment-controller    Scaled up replica set website-205870431 to 3

user@ubuntu:~/dep$

```

Checking the event log occasionally will help you identify normal cluster patterns and make it possible for you to spot anomalies more easily when debugging.

When many resources are running on a cluster it can be advantageous to restrict output to a certain set of resources. The Kubernetes labeling system makes this easy. The `-l` switch can be used with the `kubectl get` subcommand to filter output by label.

Try listing all pods:

```

user@ubuntu:~/dep$ kubectl get pods

NAME                                READY    STATUS    RESTARTS   AGE
website-205870431-7cd9k            1/1      Running   0           2m
website-205870431-h530h            1/1      Running   0           2m
website-205870431-l7czq            1/1      Running   0           2m

user@ubuntu:~/dep$

```

Now try filtering by the "appname" label key we assigned to all of our pods in the pod template metadata:

```

user@ubuntu:~/dep$ kubectl get pods -l appname

NAME                                READY    STATUS    RESTARTS   AGE
website-205870431-7cd9k            1/1      Running   0           2m
website-205870431-h530h            1/1      Running   0           2m
website-205870431-l7czq            1/1      Running   0           2m

user@ubuntu:~/dep$

```

You can also filter by key and value:

```
user@ubuntu:~/dep$ kubectl get pods -l appname=webserver
```

NAME	READY	STATUS	RESTARTS	AGE
website-205870431-7cd9k	1/1	Running	0	2m
website-205870431-h530h	1/1	Running	0	2m
website-205870431-l7czq	1/1	Running	0	2m

```
user@ubuntu:~/dep$
```

You can filter by pod name:

```
user@ubuntu:~/dep$ kubectl get $(kubectl get pods -o name | head -1)
```

NAME	READY	STATUS	RESTARTS	AGE
website-205870431-7cd9k	1/1	Running	0	3m

```
user@ubuntu:~/dep$
```

Our pod has labels we have added and the Kubernetes infrastructure may add labels as well:

```
user@ubuntu:~/dep$ kubectl describe $(kubectl get pods -o name | head -1) | grep -A2 -i label
```

Labels:

- appname=webserver
- pod-template-hash=2870940145
- targetenv=demo

```
user@ubuntu:~/dep$
```

Unfortunately describe doesn't allow for JSON output. Good news, though, `get` does.

```
user@ubuntu:~/dep$ kubectl get $(kubectl get pods -o name | head -1) -o json | jq .metadata.labels
```

```
{
  "appname": "webserver",
```

```
"pod-template-hash": "2870940145",  
"targetenv": "demo"  
}
```

```
user@ubuntu:~/dep$
```

- Why do each of the filters above work or not work?
- Enter a command to display all of the pods with either the “demo” or “prod” value for targetenv
- Find all pods other than those with the “demo” or “prod” value for targetenv
- Enter a command to display all of the pods with either the “demo” or “prod” value for targetenv and the appname key set to webserver

2. Checking status of a Deployment

We have seen previously how to check the status of a deployment.

```
user@ubuntu:~/dep$ kubectl get deploy
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
website	3	3	3	3	2m

```
user@ubuntu:~/dep$
```

Now we take an slightly more application-centric view.

```
user@ubuntu:~/dep$ kubectl rollout status deploy/website
```

```
deployment "website" successfully rolled out
```

```
user@ubuntu:~/dep$
```

Rollouts are used to update a given set of Pods, the ones controlled by this Deployment's replica set. It reports success when all the currently deployed Pods match what is expected in the current deployment. In k8s technical terms these conditions are all true:

- `.status.observedGeneration >= .metadata.generation`
- `.status.updatedReplicas == .spec.replicas`
- `.spec.availableReplicas >= minimum required`

3. Updating a Deployment

We are using nginx 1.7.9 in our example, lets update to 1.9.1.

```
user@ubuntu:~/dep$ kubectl set image deploy/website podweb=nginx:1.9.1 --record  
deployment "website" image updated  
user@ubuntu:~/dep$
```

Alternative is to use `kubectl edit deployment/website`

Check the status of the rollout:

```
user@ubuntu:~/dep$ kubectl rollout status deploy/website  
deployment "website" successfully rolled out  
user@ubuntu:~/dep$
```

```
user@ubuntu:~/dep$ kubectl get deploy/website
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
website	3	3	3	3	5m

```
user@ubuntu:~/dep$
```

Look at the Replica Sets & Pods

```
user@ubuntu:~/dep$ kubectl get rs,pod
```

NAME	DESIRED	CURRENT	READY	AGE
rs/website-205870431	0	0	0	6m


```
rs/website-4061902189    3          3          3          55s

NAME                      READY      STATUS      RESTARTS   AGE
po/website-4061902189-7kz2j 1/1        Running    0           44s
po/website-4061902189-kt956 1/1        Running    0           55s
po/website-4061902189-v11lw 1/1        Running    0           43s

user@ubuntu:~/dep$
```

By describing the deployment we can inspect the events that occurred during the rollout:

```
user@ubuntu:~/dep$ kubectl describe deploy/website

Name:                      website
Namespace:                 default
CreationTimestamp:         Tue, 12 Sep 2017 22:12:17 -0700
Labels:                    bu=sales
Annotations:               deployment.kubernetes.io/revision=2
Selector:                  appname=webserver,targetenv=demo
Replicas:                  3 desired | 3 updated | 3 total | 3 available | 0 unavailable
StrategyType:              RollingUpdate
MinReadySeconds:           0
RollingUpdateStrategy:     25% max unavailable, 25% max surge
Pod Template:
  Labels:                  appname=webserver
                           targetenv=demo
  Containers:
    podweb:
      Image:               nginx:1.9.1
      Port:                80/TCP
      Environment:         <none>
      Mounts:              <none>
      Volumes:             <none>
  Conditions:
    Type      Status  Reason
    ----      -
    Available  True   MinimumReplicasAvailable
    Progressing True   NewReplicaSetAvailable
OldReplicaSets: <none>
NewReplicaSet:  website-4061902189 (3/3 replicas created)
Events:
```

FirstSeen Message -----	LastSeen -----	Count -----	From -----	SubObjectPath -----	Type -----	Reason -----
7m	7m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled up replica set website-205870431 to 3						
1m	1m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled up replica set website-4061902189 to 1						
1m	1m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled down replica set website-205870431 to 2						
1m	1m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled up replica set website-4061902189 to 2						
1m	1m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled down replica set website-205870431 to 1						
1m	1m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled up replica set website-4061902189 to 3						
1m	1m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled down replica set website-205870431 to 0						
user@ubuntu:~/dep\$						

Note that the rollout was a smooth transition from one set of Pods controlled by our original ReplicaSet `website-205870431` to our second set of Pods controlled by the RS `website-4061902189`.

4. Manually rolling back a deployment

Lets manually revert back to nginx 1.7.9 and check the status.

```
user@ubuntu:~/dep$ kubectl set image deploy/website podweb=nginx:1.7.9 --record
deployment "website" image updated
user@ubuntu:~/dep$
```

```
user@ubuntu:~/dep$ kubectl rollout status deploy/website
deployment "website" successfully rolled out
```

```
user@ubuntu:~/dep$
```

```
user@ubuntu:~/dep$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
website-205870431	3	3	3	9m
website-4061902189	0	0	0	4m

```
user@ubuntu:~/dep$
```

Notice which deployment (NAME) is being used.

```
user@ubuntu:~/dep$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
website-205870431-4bwfs	1/1	Running	0	35s
website-205870431-62fv5	1/1	Running	0	38s
website-205870431-x65g5	1/1	Running	0	36s

```
user@ubuntu:~/dep$
```

Confirm your observations once again in the event log.

```
user@ubuntu:~/dep$ kubectl describe deploy/website
```

```
Name:                website
Namespace:           default
CreationTimestamp:    Tue, 12 Sep 2017 22:12:17 -0700
Labels:               bu=sales
Annotations:          deployment.kubernetes.io/revision=3
Selector:             appname=webserver,targetenv=demo
Replicas:             3 desired | 3 updated | 3 total | 3 available | 0 unavailable
StrategyType:         RollingUpdate
MinReadySeconds:      0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
```

```

Labels:      appname=webserver
             targetenv=demo
Containers:
  podweb:
    Image:      nginx:1.7.9
    Port:      80/TCP
    Environment: <none>
    Mounts:     <none>
  Volumes:     <none>
Conditions:
  Type            Status  Reason
  ----            -
  Available       True    MinimumReplicasAvailable
  Progressing     True    NewReplicaSetAvailable
OldReplicaSets:  <none>
NewReplicaSet:   website-205870431 (3/3 replicas created)
Events:
  FirstSeen      LastSeen        Count   From                    SubObjectPath  Type            Reason
  ----
10m             10m             1       deployment-controller   Scaled up replica set website-205870431 to 3      Normal          ScalingReplicaSet
4m              4m              1       deployment-controller   Scaled up replica set website-4061902189 to 1      Normal          ScalingReplicaSet
4m              4m              1       deployment-controller   Scaled down replica set website-205870431 to 2      Normal          ScalingReplicaSet
4m              4m              1       deployment-controller   Scaled up replica set website-4061902189 to 2      Normal          ScalingReplicaSet
4m              4m              1       deployment-controller   Scaled down replica set website-205870431 to 1      Normal          ScalingReplicaSet
4m              4m              1       deployment-controller   Scaled up replica set website-4061902189 to 3      Normal          ScalingReplicaSet
4m              4m              1       deployment-controller   Scaled down replica set website-205870431 to 0      Normal          ScalingReplicaSet
56s             56s             1       deployment-controller   Scaled up replica set website-205870431 to 1      Normal          ScalingReplicaSet
54s             54s             1       deployment-controller   Scaled down replica set website-4061902189 to 2      Normal          ScalingReplicaSet
54s             49s             4       deployment-controller   (combined from similar events): Scaled down replica set website-4061902189 to 0
user@ubuntu:~/dep$

```

5. Checking rollout history of a Deployment

We can use the *rollout history* subcommand to see what we have been doing to trigger these rollouts

```
user@ubuntu:~/dep$ kubectl rollout history deploy/website

deployments "website"
REVISION      CHANGE-CAUSE
2             kubectl set image deploy/website podweb=nginx:1.9.1 --record=true
3             kubectl set image deploy/website podweb=nginx:1.7.9 --record=true

user@ubuntu:~/dep$
```

Take a detailed look at a previous deployment version.

```
user@ubuntu:~/dep$ kubectl rollout history deploy/website --revision=2

deployments "website" with revision #2
Pod Template:
  Labels:      appname=webserver
              pod-template-hash=4061902189
              targetenv=demo
  Containers:
    podweb:
      Image:    nginx:1.9.1
      Port:    80/TCP
      Environment:  <none>
      Mounts:     <none>
      Volumes:     <none>

user@ubuntu:~/dep$
```

6. Rolling back to a previous Deployment

Confirm the current version of a container is 1.7.9.

```
user@ubuntu:~/dep$ kubectl get pods -o json | jq .items[0].spec.containers[0].image -r
nginx:1.7.9
user@ubuntu:~/dep$
```

Revert to previous version/revision.

```
user@ubuntu:~/dep$ kubectl rollout undo deploy/website
deployment "website" rolled back
user@ubuntu:~/dep$
```

Alternative to above is `kubectl rollout undo deployment/website --to-revision=2`

```
user@ubuntu:~/dep$ kubectl get deploy/website
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
website	3	3	3	3	12m

```
user@ubuntu:~/dep$
```

```
user@ubuntu:~/dep$ kubectl describe deploy/website
```

Name: website
Namespace: default
CreationTimestamp: Tue, 12 Sep 2017 22:12:17 -0700
Labels: bu=sales
Annotations: deployment.kubernetes.io/revision=4
Selector: appname=webserver,targetenv=demo
Replicas: 3 desired | 3 updated | 3 total | 3 available | 0 unavailable
StrategyType: RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge

Pod Template:

Labels: appname=webserver
targetenv=demo

Containers:

podweb:

Image: nginx:1.9.1

Port: 80/TCP

Environment: <none>

Mounts: <none>

Volumes: <none>

Conditions:

Type	Status	Reason
Available	True	MinimumReplicasAvailable
Progressing	True	NewReplicaSetAvailable

OldReplicaSets: <none>

NewReplicaSet: website-4061902189 (3/3 replicas created)

Events:

FirstSeen Message	LastSeen	Count	From	SubObjectPath	Type	Reason
12m	12m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled up replica set website-205870431 to 3						
7m	7m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled up replica set website-4061902189 to 1						
7m	7m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled down replica set website-205870431 to 2						
7m	7m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled up replica set website-4061902189 to 2						
3m	3m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled up replica set website-205870431 to 1						
3m	3m	1	deployment-controller		Normal	ScalingReplicaSet
Scaled down replica set website-4061902189 to 2						
24s	24s	1	deployment-controller		Normal	DeploymentRollback
Rolled back deployment "website" to revision 2						
3m	23s	7	deployment-controller		Normal	ScalingReplicaSet
(combined from similar events): Scaled up replica set website-4061902189 to 2						
7m	22s	2	deployment-controller		Normal	ScalingReplicaSet
Scaled up replica set website-4061902189 to 3						
7m	22s	2	deployment-controller		Normal	ScalingReplicaSet
Scaled down replica set website-205870431 to 1						
6m	20s	2	deployment-controller		Normal	ScalingReplicaSet

```
Scaled down replica set website-205870431 to 0
```

```
user@ubuntu:~/dep$
```

Note the unique event in the log for the rollback: `DeploymentRollback`

Confirm the container image version has been reverted to 1.9.1:

```
user@ubuntu:~/dep$ kubectl get pods -o json | jq .items[0].spec.containers[0].image -r
```

```
nginx:1.9.1
```

```
user@ubuntu:~/dep$
```

7. Pausing and resuming a Deployment

In a larger installation, we may be deploying dozens of pods. For our small test it is hard to pause in time, so we chain the commands to hopefully catch it in the act.

```
user@ubuntu:~/dep$ kubectl set image deploy/website podweb=nginx:1.7.9; kubectl rollout pause deploy/website
```

```
deployment "website" image updated
```

```
deployment "website" paused
```

```
user@ubuntu:~/dep$
```

```
user@ubuntu:~/dep$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
website-205870431	1	1	1	13m
website-4061902189	3	3	3	8m

```
user@ubuntu:~/dep$
```



```
user@ubuntu:~/dep$ kubectl rollout status deploy/website

Waiting for rollout to finish: 1 out of 3 new replicas have been updated...

^C

user@ubuntu:~/dep$
```

```
user@ubuntu:~/dep$ kubectl rollout resume deploy/website

deployment "website" resumed

user@ubuntu:~/dep$
```

```
user@ubuntu:~/dep$ kubectl rollout status deploy/website

deployment "website" successfully rolled out

user@ubuntu:~/dep$
```

```
user@ubuntu:~/dep$ kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
website-205870431	3	3	3	14m
website-4061902189	0	0	0	9m

```
user@ubuntu:~/dep$
```

Delete your deployment.

8. Health Checks

In this step we will create a pod with a health check. Enter and run the following config (*hc.yaml*):

```
user@ubuntu:~/dep$ cd
user@ubuntu:~$ mkdir hc
user@ubuntu:~$ cd hc
user@ubuntu:~/hc$ vi hc.yaml
user@ubuntu:~/hc$ cat hc.yaml

apiVersion: apps/v1beta1
kind: Deployment
metadata:
  name: nginx
  labels:
    name: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      name: nginx
  template:
    metadata:
      labels:
        name: nginx
    spec:
      containers:
      - name: nginx
        image: nginx:latest
        ports:
        - containerPort: 80
        livenessProbe: # An HTTP health check
          httpGet:
            path: /
            port: 80
          initialDelaySeconds: 30
          timeoutSeconds: 1

user@ubuntu:~/hc$
```

Now run the deployment:

```

user@ubuntu:~/hc$ kubectl create -f hc.yaml

deployment "nginx" created

user@ubuntu:~/hc$

```

View your deployment:

```

user@ubuntu:~/hc$ kubectl get deploy,rs,pods

```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deploy/nginx	3	3	3	3	14s

NAME	DESIRED	CURRENT	READY	AGE
rs/nginx-101291927	3	3	3	14s

NAME	READY	STATUS	RESTARTS	AGE
po/nginx-101291927-n78jd	1/1	Running	0	14s
po/nginx-101291927-n7w09	1/1	Running	0	14s
po/nginx-101291927-x38cd	1/1	Running	0	14s

```

user@ubuntu:~/hc$

```

Note that our nginx service listens on port 80 and responds normally to requests for "/", so our health check is passing.

To trigger the health check repair logic, we need to simulate an error condition. By forcing nginx to report a 404, the `HttpGet` livenessProbe will fail. We can do this by deleting the nginx configuration file in the nginx container.

Display the events for the first pod in the set:

```

user@ubuntu:~/hc$ kubectl get events | grep $(kubectl get pods -o name | head -1 | awk -F '/' '{print $2}')

```

8m	8m	1	nginx-101291927-n78jd	Pod		Normal	
Scheduled			default-scheduler	Successfully assigned nginx-101291927-n78jd to ubuntu			
8m	8m	1	nginx-101291927-n78jd	Pod	spec.containers{nginx}	Normal	Pulling
kubelet, ubuntu				pulling image "nginx:latest"			
8m	8m	1	nginx-101291927-n78jd	Pod	spec.containers{nginx}	Normal	Pulled
kubelet, ubuntu				Successfully pulled image "nginx:latest"			
8m	8m	1	nginx-101291927-n78jd	Pod	spec.containers{nginx}	Normal	Created

```

kubenet, ubuntu      Created container with id f1c3413b14d4ad1389d698846f3dd5259cefadce5a88bb6ab424f40e42309b5c
8m      8m      1      nginx-101291927-n78jd      Pod      spec.containers{nginx}      Normal      Started
kubenet, ubuntu      Started container with id f1c3413b14d4ad1389d698846f3dd5259cefadce5a88bb6ab424f40e42309b5c
8m      8m      1      nginx-101291927      ReplicaSet      Normal
SuccessfulCreate      replicaset-controller      Created pod: nginx-101291927-n78jd

user@ubuntu:~/hc$

```

The status is good.

Now lets tell the nginx in the first pod to stop serving the root IRI by deleting the nginx default config.

```

user@ubuntu:~/hc$ kubectl exec -it $(kubectl get pods -o name | head -1 | awk -F '/' '{print $2}') -- sh -c "rm
/etc/nginx/conf.d/default.conf && nginx -s reload"

2017/07/14 19:09:28 [notice] 18#18: signal process started

user@ubuntu:~/hc$

```

Now redisplay the events for the pod:

```

user@ubuntu:~/hc$ kubectl get events | grep $(kubectl get pods -o name | head -1 | awk -F '/' '{print $2}')

15m      15m      1      nginx-2181692066-9cfxt      Pod      Normal
Scheduled      default-scheduler      Successfully assigned nginx-2181692066-9cfxt to ubuntu
15m      15m      1      nginx-2181692066-9cfxt      Pod      Normal
SuccessfulMountVolume      kubelet, ubuntu      MountVolume.SetUp succeeded for volume "default-token-gl2cc"
14m      15m      2      nginx-2181692066-9cfxt      Pod      spec.containers{nginx}      Normal
Pulling      kubelet, ubuntu      pulling image "nginx:latest"
14m      15m      2      nginx-2181692066-9cfxt      Pod      spec.containers{nginx}      Normal
Pulled      kubelet, ubuntu      Successfully pulled image "nginx:latest"
14m      15m      2      nginx-2181692066-9cfxt      Pod      spec.containers{nginx}      Normal
Created      kubelet, ubuntu      Created container
14m      15m      2      nginx-2181692066-9cfxt      Pod      spec.containers{nginx}      Normal
Started      kubelet, ubuntu      Started container
14m      14m      3      nginx-2181692066-9cfxt      Pod      spec.containers{nginx}      Warning
Unhealthy      kubelet, ubuntu      Liveness probe failed: Get http://10.32.0.6:80/: dial tcp
10.32.0.6:80: getsockopt: connection refused
14m      14m      1      nginx-2181692066-9cfxt      Pod      spec.containers{nginx}      Normal

```

```

Killing kubelet, ubuntu Killing container with id docker://nginx:pod "nginx-2181692066-
9cfxt_default(0657e7d2-6bae-11e7-bbf7-000c29f5877a)" container "nginx" is unhealthy, it will be killed and re-
created.
15m      15m      1      nginx-2181692066      ReplicaSet      Normal
SuccessfulCreate      replicaset-controller      Created pod: nginx-2181692066-9cfxt

user@ubuntu:~/hc$

```

As you can see the Liveness probe is now failing. The nginx container in the pod was created, started, found unhealthy, killed, created and started again.

Remove the related resources.

```

user@ubuntu:~/jobs$ kubectl delete deploy/nginx

deployment "nginx" deleted

user@ubuntu:~/jobs$

```

9. Creating a Job

In a previous lab we saw that running a pod standalone works but without an RS the pod will not restart if it crashes. Unfortunately, if we run a batch job in a pod with an RS and the pod completes the task, the RS will start the pod again.

What if we want a pod that runs only once, however, if it or the node it is running on fails before the pod completes successfully, we want the pod to be started again until it does complete successfully. Kubernetes provides a **Job** type for this scenario.

A Job is like an RC/RS that ensures that a pod runs once to completion. Imagine we want to calculate Pi. Not twice, not half of a time, but precisely once. A job would be the perfect way to run a container that calculates Pi. Enter this sample job config to compute Pi:

```

user@ubuntu:~/hc$ cd ..

user@ubuntu:~$ mkdir jobs

user@ubuntu:~$ cd jobs/

user@ubuntu:~/jobs$ vim myjob.yaml

user@ubuntu:~/jobs$ cat myjob.yaml

```

```
apiVersion: batch/v1
kind: Job
metadata:
  name: pi
spec:
  template:
    metadata:
      name: pi
    spec:
      containers:
      - name: pi
        image: perl
        command: ["perl", "-Mbignum=bpi", "-wle", "print bpi(2000)"]
        restartPolicy: Never

user@ubuntu:~/jobs$
```

The config uses apiVersion "batch/v1". The kind of object we will create is a Job. The Job will have the name "pi", as per the metadata. The spec for our Job includes a selector which will match anything with the label "app=pi".

The template for the pod the Job we'll create must have a name pi.

The spec for the pod uses a single perl container which will run the command that computes pi. We also set the restart policy to Never.

Now try running your Job:

```
user@ubuntu:~/jobs$ kubectl create -f myjob.yaml

job "pi" created

user@ubuntu:~/jobs$
```

Examine the job:

```
user@ubuntu:~/jobs$ kubectl get deploy,rs,pods,job

NAME          DESIRED  SUCCESSFUL  AGE
jobs/pi       1        1           1m
```

```
user@ubuntu:~/jobs$
```

```
user@ubuntu:~/jobs$ kubectl describe job/pi
```

```
Name:          pi
Namespace:     default
Selector:      controller-uid=7bd22273-38da-11e7-b8ef-000c2949d6f4
Labels:        controller-uid=7bd22273-38da-11e7-b8ef-000c2949d6f4
                job-name=pi
Annotations:   <none>
Parallelism:   1
Completions:   1
Start Time:    Sun, 14 May 2017 12:21:08 -0700
Pods Statuses: 0 Running / 1 Succeeded / 0 Failed
Pod Template:
  Labels:      controller-uid=7bd22273-38da-11e7-b8ef-000c2949d6f4
                job-name=pi
```

Containers:

```
pi:
  Image:      perl
  Port:
  Command:
    perl
    -Mbignum=bpi
    -wle
    print bpi(2000)
  Environment: <none>
  Mounts:      <none>
  Volumes:     <none>
```

Events:

FirstSeen	LastSeen	Count	From	SubObjectPath	Type	Reason	
Message							
-----	-----	-----	-----	-----	-----	-----	---

1m	1m	1	job-controller		Normal	SuccessfulCreate	

Created pod: pi-fhftr
user@ubuntu:~/jobs\$

The `kubectl create` subcommand processes the job request and runs our pod. Displaying the Job description shows us the name of the pod that ran the Job.

We can now dump the logs for the pod to see the result:

```
user@ubuntu:~/jobs$ kubectl logs $(kubectl get jobs -o name)
```

```
3.1415926535897932384626433832795028841971693993751058209749445923078164062862089986280348253421170679821480865132
823066470938446095505822317253594081284811174502841027019385211055596446229489549303819644288109756659334461284756
482337867831652712019091456485669234603486104543266482133936072602491412737245870066063155881748815209209628292540
917153643678925903600113305305488204665213841469519415116094330572703657595919530921861173819326117931051185480744
623799627495673518857527248912279381830119491298336733624406566430860213949463952247371907021798609437027705392171
762931767523846748184676694051320005681271452635608277857713427577896091736371787214684409012249534301465495853710
507922796892589235420199561121290219608640344181598136297747713099605187072113499999983729780499510597317328160963
185950244594553469083026425223082533446850352619311881710100031378387528865875332083814206171776691473035982534904
287554687311595628638823537875937519577818577805321712268066130019278766111959092164201989380952572010654858632788
659361533818279682303019520353018529689957736225994138912497217752834791315155748572424541506959508295331168617278
558890750983817546374649393192550604009277016711390098488240128583616035637076601047101819429555961989467678374494
482553797747268471040475346462080466842590694912933136770289891521047521620569660240580381501935112533824300355876
402474964732639141992726042699227967823547816360093417216412199245863150302861829745557067498385054945885869269956
909272107975093029553211653449872027559602364806654991198818347977535663698074265425278625518184175746728909777727
938000816470600161452491921732172147723501414419735685481613611573525521334757418494684385233239073941433345477624
168625189835694855620992192221842725502542568876717904946016534668049886272327917860857843838279679766814541009538
837863609506800642251252051173929848960841284886269456042419652850222106611863067442786220391949450471237137869609
563643719172874677646575739624138908658326459958133904780275898
```

```
user@ubuntu:~/jobs$
```

By default, a Job is complete when one Pod runs to successful completion. You can also specify that this needs to happen multiple times by specifying Job spec key *“completions”* with a value greater than 1. You can suggest how many pods should run concurrently by setting Job spec key *“parallelism”* to the number of pods you would like to have running concurrently (the value defaults to *“completions”*.) The parallelism key is just a hint and the Job may run fewer or more concurrent pods.

Jobs are complementary to Deployments. A Deployment manages pods which are not expected to terminate (e.g. web servers,) and a Job manages pods that are expected to terminate (e.g. batch jobs.)

When you are finished exploring remove the Job:

```
user@ubuntu:~/jobs$ kubectl delete job pi
```

```
job "pi" deleted
```



```
user@ubuntu:~/jobs$
```

```
user@ubuntu:~/jobs$ kubectl get deploy,rs,pods,job
```

```
No resources found.
```

```
user@ubuntu:~/jobs$
```

Congratulations you have completed the Kubernetes Deployments and Replica Sets lab!

Copyright (c) 2013-2017 RX-M LLC, Cloud Native Consulting, all rights reserved